



ANALYSIS ON TANZANIAN WATER POINTS

BUSINESS UNDERSTANDING

OVERVIEW

Tanzania, the fifth most populous country in Africa, has experienced significant economic growth over the years, however despite considerable investment in water supply infrastructure from donor funding and the government, a significant proportion of its population remains without proper access to improved drinking water. With their Millennium Development Goals (MDGs) to halve the proportion of people that do not have access to water services by 2015, Tanzania only increased its access to improved drinking water from 54 percent to 56 percent [JMP, 2015 \(https://www.unwater.org/sites/default/files/app/uploads/2020/04/WHOUNICEF-Joint-Monitoring-Program-for-Water-Supply-and-Sanitation-JMP-%e2%80%93-2015-Update - ENG.pdf\)](https://www.unwater.org/sites/default/files/app/uploads/2020/04/WHOUNICEF-Joint-Monitoring-Program-for-Water-Supply-and-Sanitation-JMP-%e2%80%93-2015-Update-ENG.pdf). The country now faces a difficult task of meeting the Sustainable Development Goals (SDGs) to provide universal coverage of safe water by 2030.

Despite their efforts one persistent problem that has adversely affected the country's effort in increasing access to improved water services is the prevailing high levels of non-functionality or failures of its current water infrastructures and in particular, water points and while this issue is prevalent in Africa, evidence indicates that the problem of water point failures may be relatively more serious in Tanzania with some estimates putting the figure as high as 44 percent [Banks & Furey, 2016](#)
(https://www.researchgate.net/publication/312027512_What's_Working_Where_and_for_How_Lo

A holistic approach was used to determine the factors to be considered to be able to predict water point failure. These factors included age of the water point, technology used, the quantity and quality of the water as well as location and management of these water points. We also considered the population that use these water points, and the sources of the water. Using a variety of statistical methods, we seek to understand the impact of these factors on water point's failure and develop a model to predict the possibility of water pump failure with an accuracy of 80% and f1 score of 75%

PROBLEM STATEMENT

We have been tasked by World Bank Group together with the Government of Tanzania to seek a better understanding as to why water point failure is significantly high in Tanzania as well as a way to reliably predict when water pumps shall fail as they tackle the difficult task of meeting their 2030 MDG goals in Environmental Sustainability .

OBJECTIVES

The research seeks to meet the following objectives:

1. Analyze the Impact of Age, Technology, and Investment on Water Point Failure

2. Assess the Impact of Socioeconomic and Geographical Factors

3. Develop a Predictive Model for Water Point Failure

DATA UNDERSTANDING

This research utilized data from [DRIVEN DATA](#)
(<https://www.drivendata.org/competitions/7/pump-it-up-data-mining-the-water-table/page/23/>) about waterpoints. The dataset was split into three CSV files:

- **Training set values**
- **Training set labels**
- **Test set values**

The training and test datasets contained similar columns, while the training set labels dataset included one column, which was the focus of the study.

Column Information

The following columns were provided in the training and testing datasets:

- **amount_tsh**: Total static head (amount of water available to waterpoint)
- **date_recorded**: The date the row was entered
- **funder**: Who funded the well
- **gps_height**: Altitude of the well
- **installer**: Organization that installed the well
- **longitude**: GPS coordinate
- **latitude**: GPS coordinate
- **wpt_name**: Name of the waterpoint if there is one
- **num_private**: [Missing information]
- **basin**: Geographic water basin
- **subvillage**: Geographic location
- **region**: Geographic location
- **region_code**: Geographic location (coded)
- **district_code**: Geographic location (coded)
- **lga**: Geographic location
- **ward**: Geographic location
- **population**: Population around the well
- **public_meeting**: True/False
- **recorded_by**: Group entering this row of data
- **scheme_management**: Who operates the waterpoint
- **scheme_name**: Who operates the waterpoint
- **permit**: If the waterpoint is permitted
- **construction_year**: Year the waterpoint was constructed
- **extraction_type**: The kind of extraction the waterpoint uses
- **extraction_type_group**: The kind of extraction the waterpoint uses
- **extraction_type_class**: The kind of extraction the waterpoint uses
- **management**: How the waterpoint is managed
- **management_group**: How the waterpoint is managed
- **payment**: What the water costs
- **payment_type**: What the water costs
- **water_quality**: The quality of the water
- **quality_group**: The quality of the water
- **quantity**: The quantity of water
- **quantity_group**: The quantity of water
- **source**: The source of the water
- **source_type**: The source of the water
- **source_class**: The source of the water
- **waterpoint_type**: The kind of waterpoint

- **waterpoint_type_group**: The kind of waterpoint

Labels Information

The labels in the training set labels contained one column, **status_group**. This column indicates the condition of the waterpoint with the following possible outcomes:

- **functional**: The waterpoint is operational and there are no repairs needed
- **functional needs repair**: The waterpoint is operational, but needs repairs
- **non functional**: The waterpoint is not operational

DATA PREPARATION

The following steps in summary shall be followed in the data preparation stage in preparation for Modeling in later stages

1. Data Loading

- Load the Datasets
- Inspect the Data

2. Data Cleaning

- Validity Check
- Consistency Check
- Uniformity Check
- Completeness Check

3. EXPLORATORY DATA ANALYSIS

- Understand Data Distribution
- Identify Relationships - Univariate and Bivariate Analysis
- Handle High Cardinality Columns

DATA LOADING

The following was carried out

1. Loading the Datasets
2. Inspecting the Data

```
In [1]: import os
import numpy as np
import pandas as pd
from itertools import combinations
import warnings
warnings.filterwarnings('ignore')

# Libraries for visualizations
import folium
import seaborn as sns
import plotly.express as px
from IPython.display import display, HTML
import matplotlib.pyplot as plt
%matplotlib inline

# Libraries for Model Preprocessing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder, LabelEncoder
from imblearn.over_sampling import SMOTE

# Libraries for Modeling
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import GradientBoostingClassifier, RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.model_selection import RandomizedSearchCV, cross_val_score
from sklearn.metrics import accuracy_score, f1_score, make_scorer, confusion_matrix
```


In [2]:

```

class DataLoader:
    def __init__(self):
        pass

    def read_data(self, file_path):
        _, file_ext = os.path.splitext(file_path)
        """
        Load data from a CSV, TSV, JSON or Excel file
        """
        if file_ext == '.csv':
            return pd.read_csv(file_path, index_col=None)

        elif file_ext == '.tsv':
            return pd.read_csv(file_path, sep='\t')

        elif file_ext == '.json':
            return pd.read_json(file_path)

        elif file_ext in ['.xls', '.xlsx']:
            return pd.read_excel(file_path)

        else:
            raise ValueError(f"Unsupported file format:")

class DataFrameMerger:
    def __init__(self):
        pass

    def merge_dataframes(self, df1, df2, on, how='inner'):
        """
        Merge two dataframes on specified columns with the specified method.
        """
        merged_df = pd.merge(df1, df2, on=on, how=how)
        return merged_df.sort_values(by=merged_df.columns.tolist())

class DataInfo:

    def __init__(self, df):
        self.df = df

    def info(self):
        """
        Displaying Relevant Information on the the Dataset Provided
        """
        # Counting no of rows
        print(f'\nTotal Rows : {self.df.shape[0]} \n' + '--'*10 )

        # Counting no of columns
        print(f'\nTotal Columns : {self.df.shape[1]} \n' + '--'*10)

        # Extracting column names
        column_name = self.df.columns
        print(f'\nColumn Names\n' + '--'*10 + f'\n{column_name} \n \n')

        # Data type info
        print(f'Data Summary\n' + '--'*10)

```

```
data_summary = self.df.info()

# Total null values by each categories
null_values = self.df.isnull().sum()
print(f'\nNull values\n' + '--'*10 + f'\n{null_values} \n \n')

# Descriptive statistics
describe = self.df.describe()
print(f'\nDescriptive Statistics\n' + '--'*10 )
display(describe)

#Display the dataset
print(f'\nDataset Overview\n'+ '--'*10)
return self.df.head()
```



```
In [3]: #Instantiate the Loader class
data_loader = DataLoader()

# Loading the datasets
train_data=data_loader.read_data("data/train.csv")
train_labels=data_loader.read_data("data/train-labels.csv")
test_data=data_loader.read_data("data/test.csv")

# Instantiate the DF merger class
merger=DataFrameMerger()

# Merge the train data provided
data=merger.merge_dataframes(train_data, train_labels, on="id")

print(f'\nTimeLine of Recorded Data\n' + '--'*10 )
print(f"From:",data['date_recorded'].min(), "To:", data['date_recorded'].max())
print(f'--'*10 )

# Instantiate the Information class
information=DataInfo(data)

# Getting the info on the training DF
information.info()
```

TimeLine of Recorded Data

 From: 2002-10-14 To: 2013-12-03

Total Rows : 59400

Total Columns : 41

Column Names

 Index(['id', 'amount_tsh', 'date_recorded', 'funder', 'gps_height',
 'installer', 'longitude', 'latitude', 'wpt_name', 'num_private',
 'basin', 'subvillage', 'region', 'region_code', 'district_code', 'lga',
 'ward', 'population', 'public_meeting', 'recorded_by',
 'scheme_management', 'scheme_name', 'permit', 'construction_year',
 'extraction_type', 'extraction_type_group', 'extraction_type_class',
 'management', 'management_group', 'payment', 'payment_type',
 'water_quality', 'quality_group', 'quantity', 'quantity_group',
 'source', 'source_type', 'source_class', 'waterpoint_type',
 'waterpoint_type_group', 'status_group'],
 dtype='object')

Data Summary

 <class 'pandas.core.frame.DataFrame'>
 Int64Index: 59400 entries, 9410 to 39131
 Data columns (total 41 columns):

#	Column	Non-Null Count	Dtype
0	id	59400 non-null	int64
1	amount_tsh	59400 non-null	float64
2	date_recorded	59400 non-null	object
3	funder	55765 non-null	object
4	gps_height	59400 non-null	int64
5	installer	55745 non-null	object
6	longitude	59400 non-null	float64
7	latitude	59400 non-null	float64
8	wpt_name	59400 non-null	object
9	num_private	59400 non-null	int64
10	basin	59400 non-null	object
11	subvillage	59029 non-null	object
12	region	59400 non-null	object
13	region_code	59400 non-null	int64
14	district_code	59400 non-null	int64
15	lga	59400 non-null	object
16	ward	59400 non-null	object
17	population	59400 non-null	int64
18	public_meeting	56066 non-null	object
19	recorded_by	59400 non-null	object
20	scheme_management	55523 non-null	object
21	scheme_name	31234 non-null	object

```

22 permit 56344 non-null object
23 construction_year 59400 non-null int64
24 extraction_type 59400 non-null object
25 extraction_type_group 59400 non-null object
26 extraction_type_class 59400 non-null object
27 management 59400 non-null object
28 management_group 59400 non-null object
29 payment 59400 non-null object
30 payment_type 59400 non-null object
31 water_quality 59400 non-null object
32 quality_group 59400 non-null object
33 quantity 59400 non-null object
34 quantity_group 59400 non-null object
35 source 59400 non-null object
36 source_type 59400 non-null object
37 source_class 59400 non-null object
38 waterpoint_type 59400 non-null object
39 waterpoint_type_group 59400 non-null object
40 status_group 59400 non-null object
dtypes: float64(3), int64(7), object(31)
memory usage: 19.0+ MB

```

Null values

```

-----
id 0
amount_tsh 0
date_recorded 0
funder 3635
gps_height 0
installer 3655
longitude 0
latitude 0
wpt_name 0
num_private 0
basin 0
subvillage 371
region 0
region_code 0
district_code 0
lga 0
ward 0
population 0
public_meeting 3334
recorded_by 0
scheme_management 3877
scheme_name 28166
permit 3056
construction_year 0
extraction_type 0
extraction_type_group 0
extraction_type_class 0
management 0
management_group 0
payment 0
payment_type 0
water_quality 0
quality_group 0

```

```
quantity          0
quantity_group    0
source            0
source_type       0
source_class      0
waterpoint_type   0
waterpoint_type_group 0
status_group      0
dtype: int64
```

Descriptive Statistics

	id	amount_tsh	gps_height	longitude	latitude	num_private	r
count	59400.000000	59400.000000	59400.000000	59400.000000	5.940000e+04	59400.000000	59
mean	37115.131768	317.650385	668.297239	34.077427	-5.706033e+00	0.474141	
std	21453.128371	2997.574558	693.116350	6.567432	2.946019e+00	12.236230	
min	0.000000	0.000000	-90.000000	0.000000	-1.164944e+01	0.000000	
25%	18519.750000	0.000000	0.000000	33.090347	-8.540621e+00	0.000000	
50%	37061.500000	0.000000	369.000000	34.908743	-5.021597e+00	0.000000	
75%	55656.500000	20.000000	1319.250000	37.178387	-3.326156e+00	0.000000	
max	74247.000000	350000.000000	2770.000000	40.345193	-2.000000e-08	1776.000000	

Dataset Overview

Out[3]:

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_
9410	0	0.0	2012-11-13	Tasaf	0	TASAF	33.125828	-5.118154	M
18428	1	0.0	2011-03-05	Shipo	1978	SHIPO	34.770717	-9.395642	
12119	2	0.0	2011-03-27	Lvia	0	LVIA	36.115056	-6.279268	Bc
10629	3	10.0	2013-06-03	Germany Republi	1639	CES	37.147432	-3.187555	Na
2343	4	0.0	2011-03-22	Cmsr	0	CMSR	36.164893	-6.099289	E:

5 rows x 41 columns

Initial Observations:

- There seem to be unique identifiers in the dataset that need to be investigated eg "id"
- Some columns seem to carry repeated data based on data understanding and column names and should be investigated
- Numeric features seem to be mainly location identifiers

DATA CLEANING

Data cleaning shall be carried out in the following steps:

1. Validity Check
2. Consistency Check
3. Uniformity Check
4. Completeness Check

VALIDITY CHECK

- Check for unique identifiers i.e. distinct elements in a column that are equal to the length of the column value counts
- Check for constant columns i.e. total distinct elements in a column is equal to 1

```
In [4]: class Validity:
    def __init__(self,df):
        self.df=df
        self.unique_identifier = []

    def find_unique_identifiers(self):
        """
        Identify unique identifiers and constant columns and drop them
        """
        self.df= self.df.copy()
        columns = self.df.columns

        for column in columns:
            if self.df[column].nunique() == len(self.df[column]):
                self.unique_identifier.append(column)
                print(f" Unique Identifier Columns:", column)
            elif self.df[column].nunique() == 1:
                self.unique_identifier.append(column)
                print(f" Constant Columns:", column)

        self.df = self.df.drop(columns=self.unique_identifier)
        display(self.df)

    return self.df
```

```
In [5]: # Checking for Validity on dataset


# Instantiate the Validity check class
valid= Validity(data)

# Validate the training dataset
train_data=valid.find_unique_identifiers()
```

Unique Identifier Columns: id
Constant Columns: recorded_by

	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name
9410	0.0	2012-11-13	Tasaf	0	TASAF	33.125828	-5.118154	Mra
18428	0.0	2011-03-05	Shipo	1978	SHIPO	34.770717	-9.395642	no
12119	0.0	2011-03-27	Lvia	0	LVIA	36.115056	-6.279268	Bomb
10629	10.0	2013-06-03	Germany Republi	1639	CES	37.147432	-3.187555	Are Namb
2343	0.0	2011-03-22	Cmsr	0	CMSR	36.164893	-6.099289	Ezelo
...	
15137	0.0	2013-03-22	World Vision	1183	World vision	37.007726	-3.280868	Upei Prim Sch
8667	0.0	2011-04-12	Danida	0	DANIDA	33.724987	-8.940758	k Mvu
22584	0.0	2012-11-13	Ministry Of Water	1188	Hesawa	33.963539	-1.429477	k Wamb Mse
108	50.0	2011-03-07	Ruthe	1428	Ruthe	35.630481	-7.710549	no
39131	50.0	2013-02-16	Mission	965	DWE	35.432998	-10.639270	k Mapui

59400 rows × 39 columns



Observations:

- 1 column had unique identifiers i.e. "id" column
- 1 column had constant column identifiers i.e. "recorded_by" column
- Dataset now has 59,400 rows and 39 columns

CONSISTENCY CHECK

- Duplicates checked in rows and dropped them
- Duplicates Checked in columns and dropped them


```

In [6]: class Consistency:
    def __init__(self,df):
        self.df=df
        self.to_drop = []

    def duplicated_rows(self):
        """
        Displaying the duplicated rows for visual assesment
        """
        df_sorted = self.df.sort_values(by=self.df.columns.tolist())

        # Find duplicated rows
        duplicates = df_sorted[df_sorted.duplicated(keep=False)]

        # Display the duplicated rows as HTML
        return display(HTML(duplicates.to_html()))

    def drop_duplicated_rows(self,rows=None):
        """
        Dropping confirmed duplicated rows
        """
        self.df.drop_duplicates(subset=rows, keep= "first", inplace= True)
        display(self.df.shape)
        return self.df

    def find_duplicated_columns(self):
        """
        Displaying the duplicated columns for visual assesment
        """
        duplicated_columns = []
        columns = self.df.columns
        for i, col1 in enumerate(columns):
            for col2 in columns[i + 1:]:
                if self.df[col1].equals(self.df[col2]):
                    duplicated_columns.append((col1, col2))
        for pair in duplicated_columns:
            display(pair[0],(self.df[pair[0]].value_counts()))
            display(pair[1],(self.df[pair[1]].value_counts()))
            self.to_drop.append(pair[1])
        return display(f"Duplicated Columns:",duplicated_columns)

    def drop_duplicate_columns(self, columns=None):
        """
        Dropping confirmed duplicated columns
        """
        if columns is None:
            columns = self.to_drop
        self.df = self.df.drop(columns=columns)
        display(self.df.shape)
        return self.df

```

```
In [7]: # Checking for Consistency in dataset

# Instantiate the Consistency check class
const= Consistency(train_data)

# Visually Checking for duplicated rows
duplicates= const.duplicated_rows()

# Dropping the duplicated rows
train_data=const.drop_duplicated_rows()
```

34310	0.0	2011-07-19	Plan International	0	Plan International	0.00000	-2.000000e
13355	0.0	2011-07-19	Plan International	0	Plan International	0.00000	-2.000000e
37202	0.0	2011-07-26	Hesawa	0	DWE	0.00000	-2.000000e
8460	0.0	2011-07-26	Hesawa	0	DWE	0.00000	-2.000000e
25300	0.0	2011-07-27	Hesawa	0	DWE	0.00000	-2.000000e
31558	0.0	2011-07-27	Hesawa	0	DWE	0.00000	-2.000000e
7907	0.0	2011-07-27	Hesawa	0	DWE	0.00000	-2.000000e
51183	0.0	2011-07-28	Government Of Tanzania	0	Government	0.00000	-2.000000e

Observations:

- 36 rows are confirmed duplicates and have been dropped
- Dataset now has 59,364 rows and 39 columns

In [8]: *# Checking for Consistency in columns dataset*

```
# Visually Checking for duplicated columns  
const.find_duplicated_columns()
```

```
# Dropping the duplicated columns  
train_data=const.drop_duplicate_columns()
```

```
'quantity'
```

```
enough          33165  
insufficient     15119  
dry              6243  
seasonal         4048  
unknown          789  
Name: quantity, dtype: int64
```

```
'quantity_group'
```

```
enough          33165  
insufficient     15119  
dry              6243  
seasonal         4048  
unknown          789  
Name: quantity_group, dtype: int64
```

```
'Duplicated Columns:'
```

```
[('quantity', 'quantity_group')]
```

```
(59364, 38)
```

Observations:

- 1 column visually confirmed to be a duplicate i.e. "quantity_group" and dropped
- Dataset now has 59,364 rows and 38 columns

UNIFORMITY CHECK

- Assessing Data Distributions i.e. outliers
- Checking Data Types


```

In [9]: class Uniformity:
    def __init__(self, df):
        self.df = df
        self.categorical_columns = []
        self.numerical_columns = []

    def column_seperation(self):
        """
        Seperate the columns into Categorical and Numerical Columns
        """
        for col in self.df.columns:
            if self.df[col].dtype == object:
                self.categorical_columns.append(col)
            else:
                self.numerical_columns.append(col)

        return self.numerical_columns, self.categorical_columns

    def detect_outliers_iqr(self):
        """
        Detect outliers in numerical columns using the IQR method.
        """
        outlier_columns = []

        for column in self.numerical_columns:
            Q1 = self.df[column].quantile(0.25)
            Q3 = self.df[column].quantile(0.75)
            IQR = Q3 - Q1
            lower_bound = Q1 - 1.5 * IQR
            upper_bound = Q3 + 1.5 * IQR

            outlier_indices = self.df[(self.df[column] < lower_bound) | (self.df[column] > upper_bound)]

            if outlier_indices:
                outlier_columns.append(column)

        return outlier_columns

    def plot_outliers(self, outlier_columns):
        """
        Plot boxplots for the columns that have outliers using Seaborn.
        """
        num_rows = (len(outlier_columns) + 2) // 3
        num_cols = min(len(outlier_columns), 3)
        fig, axes = plt.subplots(nrows=num_rows, ncols=num_cols, figsize=(20, 8))
        axes = axes.flatten()

        for i, column in enumerate(outlier_columns):
            sns.boxplot(x=self.df[column], ax=axes[i])
            axes[i].set_xlabel(column)
            axes[i].set_ylabel('Values')
            axes[i].set_title(f'{column}')
            axes[i].tick_params(axis='x', rotation=45)

        # Adjust Layout to prevent overLapping
        plt.tight_layout()

        # Show the plots

```

```
plt.show()
```

```
def convert_column_dtype(self, column_name, dtype):
    """
    Convert the data type of a column in a DataFrame.
    """
    self.df[column_name] = self.df[column_name].astype(dtype).astype(dtype)

    return self.df
```

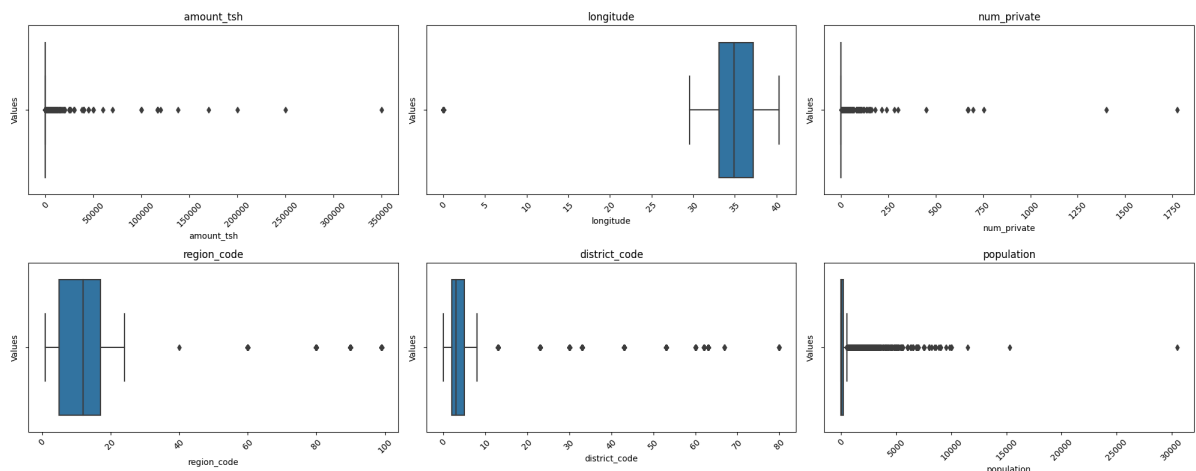
```
In [10]: #Instantiate the Uniformity class
uniform = Uniformity(train_data)

# Separating the columns into numerical and categorical
numerical_columns, categorical_columns=uniform.column_seperation()
print(f"Numerical Columns:",numerical_columns)

# Looking for Outliers in the numerical_columns
outlier_columns = uniform.detect_outliers_iqr()
print("Columns with outliers detected using IQR method:", outlier_columns)

# Plotting the columns with outliers
uniform.plot_outliers(outlier_columns)
```

Numerical Columns: ['amount_tsh', 'gps_height', 'longitude', 'latitude', 'num_private', 'region_code', 'district_code', 'population', 'construction_year']
Columns with outliers detected using IQR method: ['amount_tsh', 'longitude', 'num_private', 'region_code', 'district_code', 'population']



Observation:

- Region_code and District code are geographical code identifiers and shall be dropped since these codes regularly change and are not the current one based on their associated region columns [reference](#)
([https://www.iso.org/obp/ui/#iso:code:3166:TZ](#))

```
In [11]: #Further analysis of the Longitude and Latitude Columns
# check to see if latitude outlier is the same for both regions
mwanza=train_data[train_data["region"]=="Mwanza"]["latitude"].max()
shinyanga=train_data[train_data["region"]=="Shinyanga"]["latitude"].max()
mwanza==shinyanga
```

Out[11]: True

```
In [12]: # Reassigning it as maxim
maxim =mwanza

# Check the outlier value in maxim variable for Longitude based on the region c
train_data[train_data["latitude"]==maxim]["region"].value_counts()
```

```
Out[12]: Shinyanga    1003
Mwanza              774
Name: region, dtype: int64
```

```
In [13]: # Check the outlier value 0 for Longitude based on the region column
train_data[train_data.longitude == 0]["region"].value_counts()
```

```
Out[13]: Shinyanga    1003
Mwanza              774
Name: region, dtype: int64
```

```
In [14]: # Function to replace all outlier values with the median for the Longitude column
def replace_zero_longitudes(df, longitude_col='longitude', latitude_col='latitude', region_col='region'):
    """
    Replace longitude and Latitude Outlier values with the median longitude of each region
    """
    regions = df[region_col].unique()
    median_longitudes = {}
    median_latitudes = {}

    for region in regions:
        median_longitude = df[df[region_col] == region][longitude_col].median()
        median_latitude = df[df[region_col] == region][latitude_col].median()
        median_longitudes[region] = median_longitude
        median_latitudes[region] = median_latitude

    df.loc[(df[region_col] == region) & (df[longitude_col] == 0), longitude_col] = median_longitudes[region]
    df.loc[(df[region_col] == region) & (df[latitude_col] == maxim), latitude_col] = median_latitudes[region]

    return df, median_longitudes, median_latitudes

train_data, median_longitudes, median_latitudes = replace_zero_longitudes(train_data)
print("\nMedian Longitude values used:")
print("-----")
print("Shinyanga:", median_longitudes["Shinyanga"])
print("Mwanza:   ", median_longitudes["Mwanza"])
print("\nMedian Latitude values used:")
print("-----")
print("Shinyanga:", median_latitudes["Shinyanga"])
print("Mwanza:   ", median_latitudes["Mwanza"])
train_data
```

Median Longitude values used:

Shinyanga: 33.04787608

Mwanza: 32.99787276

Median Latitude values used:


Shinyanga: -3.3524871149999997

Mwanza: -2.52297834

Out[14]:

	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name
9410	0.0	2012-11-13	Tasaf	0	TASAF	33.125828	-5.118154	Mrai
18428	0.0	2011-03-05	Shipo	1978	SHIPO	34.770717	-9.395642	no
12119	0.0	2011-03-27	Lvia	0	LVIA	36.115056	-6.279268	Bomb
10629	10.0	2013-06-03	Germany Republi	1639	CES	37.147432	-3.187555	Are Namb
2343	0.0	2011-03-22	Cmsr	0	CMSR	36.164893	-6.099289	Ezela
...
15137	0.0	2013-03-22	World Vision	1183	World vision	37.007726	-3.280868	Upei Prim Sch
8667	0.0	2011-04-12	Danida	0	DANIDA	33.724987	-8.940758	k Mvu
22584	0.0	2012-11-13	Ministry Of Water	1188	Hesawa	33.963539	-1.429477	k Wamb Msa
108	50.0	2011-03-07	Ruthe	1428	Ruthe	35.630481	-7.710549	no
39131	50.0	2013-02-16	Mission	965	DWE	35.432998	-10.639270	k Mapui

59364 rows x 38 columns

**Observations:**

- It was evident that the longitude and latitude columns for Mwanza and Shinyanga had been given incorrect values i.e. outlier positions.
- The longitudes and latitudes have been replaced by the median of the regions and is a suitable replacement and can be confirmed from [TZ_Cities_db \(https://simplemaps.com/data/tz-cities\)](https://simplemaps.com/data/tz-cities)
- Median was used as it is least affected by the outlier values. It would also give the most probable accurate value
- Dataset has 59,364 rows and 38 columns

In [15]: *# Convert the list datatype to bool and return info on the data*

```
list_to_bool=["permit", "public_meeting"]
uniform.convert_column_dtype(list_to_bool, bool).info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 59364 entries, 9410 to 39131
Data columns (total 38 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   amount_tsh                           59364 non-null  float64
1   date_recorded                         59364 non-null  object
2   funder                               55729 non-null  object
3   gps_height                           59364 non-null  int64
4   installer                           55709 non-null  object
5   longitude                           59364 non-null  float64
6   latitude                            59364 non-null  float64
7   wpt_name                             59364 non-null  object
8   num_private                          59364 non-null  int64
9   basin                               59364 non-null  object
10  subvillage                          58993 non-null  object
11  region                              59364 non-null  object
12  region_code                         59364 non-null  int64
13  district_code                      59364 non-null  int64
14  lga                                 59364 non-null  object
15  ward                               59364 non-null  object
16  population                          59364 non-null  int64
17  public_meeting                      59364 non-null  bool
18  scheme_management                  55487 non-null  object
19  scheme_name                        31225 non-null  object
20  permit                             59364 non-null  bool
21  construction_year                  59364 non-null  int64
22  extraction_type                    59364 non-null  object
23  extraction_type_group              59364 non-null  object
24  extraction_type_class              59364 non-null  object
25  management                        59364 non-null  object
26  management_group                  59364 non-null  object
27  payment                           59364 non-null  object
28  payment_type                      59364 non-null  object
29  water_quality                     59364 non-null  object
30  quality_group                     59364 non-null  object
31  quantity                         59364 non-null  object
32  source                           59364 non-null  object
33  source_type                       59364 non-null  object
34  source_class                      59364 non-null  object
35  waterpoint_type                   59364 non-null  object
36  waterpoint_type_group             59364 non-null  object
37  status_group                      59364 non-null  object
dtypes: bool(2), float64(3), int64(6), object(27)
memory usage: 16.9+ MB
```

Observation:

- Permit and public meeting have now been accurately classified as boolean types
- All other datatypes shall be used as is

COMPLETENESS CHECK

- Check for similarity in columns
- Check for Null values in all rows


```

In [16]: class Completeness:
    def __init__(self,df):
        self.df=df

    def similarity(self,threshold=0.2):
        """
        Identifying columns that have similarity in their value counts
        with a low threshold to avoid missing similarities
        """
        similar_columns = []

        # Calculate the similarity between each pair of columns
        for col1, col2 in combinations(self.df.columns, 2):
            set1, set2 = set(self.df[col1]), set(self.df[col2])
            intersection = len(set1 & set2)
            union = len(set1 | set2)
            similarity = intersection / union if union != 0 else 0

            # Check if similarity exceeds threshold
            if similarity > threshold:
                similar_columns.append((col1, col2))

        return similar_columns

    def similar_columns(self,similar_columns):
        """
        Visually representing these similar columns for inspection
        """
        for col1, col2 in similar_columns:
            print(f"Value counts for columns '{col1}' and '{col2}':")
            print("\nColumn '{col1}' value counts:".format(col1))
            print(self.df[col1].value_counts())
            print("\nColumn '{col2}' value counts:".format(col2))
            print(self.df[col2].value_counts())
            print("\n")

    def null_values(self):
        """
        Identify Null values in dataset as value count and percentage
        """
        # Get features with null values
        null_features = self.df.columns[self.df.isnull().any()].tolist()

        # Calculate the number of missing values for each feature
        null_counts = self.df[null_features].isnull().sum()

        # Calculate the percentage of missing data for each feature
        null_percentages = self.df[null_features].isnull().mean() * 100

        # Create a DataFrame to display the results
        null_info = pd.DataFrame({
            'Column Names': null_features,
            'Missing Values': null_counts,
            'Percentage Missing': null_percentages
        }).reset_index(drop=True)

        return null_info

```

```

def handle_missing_values(self):
    """
    Handle missing values in the DataFrame.
    """
    null_info = self.null_values()

    # Apply conditions for handling missing values
    for index, row in null_info.iterrows():
        if row['Percentage Missing'] < 5:
            # Drop rows with missing values
            self.df.dropna(subset=[row['Column Names']], inplace=True)
        elif 5 <= row['Percentage Missing'] <= 10:
            # Replace missing values with "Unknown"
            self.df[row['Column Names']].fillna("Unknown", inplace=True)
        else:
            # More than 10% missing, mark column for dropping
            print(f"Column '{row['Column Names']}' has more than 10% missing")

    return self.df

```

In [17]: # Checking for completeness in the dataset

```

# Instantiate the Completeness check class
comp= Completeness(train_data)

# Find columns that have similar distinct elements using a small threshold of 0.2
similar_columns=comp.similarity(threshold=0.2)
display(similar_columns)

# Visually comparing similar columns
comp.similar_columns(similar_columns)

```

```

[('gps_height', 'population'),
 ('num_private', 'region_code'),
 ('region_code', 'district_code'),
 ('public_meeting', 'permit'),
 ('extraction_type', 'extraction_type_group'),
 ('extraction_type_group', 'extraction_type_class'),
 ('management', 'management_group'),
 ('payment', 'payment_type'),
 ('water_quality', 'quality_group'),
 ('source', 'source_type'),
 ('waterpoint_type', 'waterpoint_type_group')]

```

Value counts for columns 'gps_height' and 'population':

Column 'gps_height' value counts:

```

0      20402
-15      60
-13      55
-16      55
1790      57

```

Observations:

- After visually inspecting the column pairs, some columns bare similarity (i.e. a column is a subset of another column) to each other and were dropped to avoid repeated data points.
- The columns are:

```
'waterpoint_type_group', 'source', 'water
_quality', 'payment',
      'management', 'extraction_type_group',
'extraction_type'
```

```
In [18]: # Inspecting for null values
comp.null_values()
```

Out[18]:

	Column Names	Missing Values	Percentage Missing
0	funder	3635	6.123240
1	installer	3655	6.156930
2	subvillage	371	0.624958
3	scheme_management	3877	6.530894
4	scheme_name	28139	47.400782


```
In [19]: # Handling with null values
comp.handle_missing_values()
```

Column 'scheme_name' has more than 10% missing values.

Out[19]:

	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name
9410	0.0	2012-11-13	Tasaf	0	TASAF	33.125828	-5.118154	Mrai
18428	0.0	2011-03-05	Shipo	1978	SHIPO	34.770717	-9.395642	no
12119	0.0	2011-03-27	Lvia	0	LVIA	36.115056	-6.279268	Bomb
10629	10.0	2013-06-03	Germany Republi	1639	CES	37.147432	-3.187555	Are Namb
2343	0.0	2011-03-22	Cmsr	0	CMSR	36.164893	-6.099289	Ezek
...
15137	0.0	2013-03-22	World Vision	1183	World vision	37.007726	-3.280868	Upe Prim Sch
8667	0.0	2011-04-12	Danida	0	DANIDA	33.724987	-8.940758	k Mvu
22584	0.0	2012-11-13	Ministry Of Water	1188	Hesawa	33.963539	-1.429477	k Wamb Mse
108	50.0	2011-03-07	Ruthe	1428	Ruthe	35.630481	-7.710549	no
39131	50.0	2013-02-16	Mission	965	DWE	35.432998	-10.639270	k Mapui

58993 rows x 38 columns



Observation:

- The null values have been handled according to their percentage of missing values

- Less than 5% values have had those rows dropped i.e. subvillage column(371 rows)
- Between 5-10% have had their Nan Values replaced by Unknown i.e. scheme_management, funder and installer columns
- Greater than 10% will be added to a features drop list to be dropped i.e, scheme_name column


```
In [20]: # After visual inspection, generating list of similar columns to drop
features_to_drop=['waterpoint_type', 'source', 'water_quality', 'payment',
                  'management', 'extraction_type_group', 'extraction_type',
                  'scheme_name', "region_code", "district_code", "num_private",
                  ]

# recalling drop duplicates from the consistency class
clean_data=const.drop_duplicate_columns(features_to_drop)

(58993, 27)
```

Observation:

- The dataset now has 58,993 columns and 27 rows
- The dataset still has columns with very high cardinality. This shall be taken care of before modelling as they are still important for bivariate analysis

EXPLORATORY DATA ANALYSIS

In this section consideration was given to:

- Feature Engineering
- Univariate Analysis
- Bivariate Analysis
- Handling High Cardinality Columns

FEATURE ENGINEERING

Before continuing with the EDA process, the following needs to be engineered

- Target variable shall be converted to a binary classification problem.
- Age of the water points is of vital importance and so feature engineering shall also be carried out to ascertain the age of these points and only those with a viable age were used

```
In [21]: # Define the names to be changed and their new values
name_changes = {
    "functional needs repair": 'functional',
}

# Replace the names in the status_group column
clean_data['status_group'] = clean_data['status_group'].replace(name_changes)
clean_data.status_group.value_counts()
```

```
Out[21]: functional      36345
non functional    22648
Name: status_group, dtype: int64
```

```
In [22]: # Function to calculate the age of the water points
def calculate_age(df, date_col='date_recorded', year_col='construction_year',
    """
    Processes the date column to extract the year and calculates the age based
    """
    # Extract the year from the date column and replace the column with the year
    df[date_col] = pd.to_datetime(df[date_col]).dt.year

    # Calculate the age and create a new column
    df[new_col] = df[date_col]-df[year_col]

    return df

# Process the DataFrame and calculate age
clean_data = calculate_age(clean_data)

#Filter the dataset to only have data with relevant age
aged_data = clean_data[(clean_data["age"] >= 0) & (clean_data["age"] <= 100)]
print(f"Shape of the cleaned_data:",clean_data.shape)
print(f"Shape of the aged_data:",aged_data.shape)

Shape of the cleaned_data: (58993, 28)
Shape of the aged_data: (38672, 28)
```

```
In [23]: # Value_counts to see distribution of data without construction year
deleted_data = clean_data[ (clean_data["age"] >= 100)]
deleted_data["region"].value_counts()
```

```
Out[23]: Shinyanga      4816
Mbeya      4639
Kagera     3315
Mwanza     2714
Tabora     1959
Dodoma     1840
Iringa     372
Pwani      176
Lindi      113
Tanga      94
Mtwara     87
Arusha     51
Kilimanjaro 35
Morogoro   18
Dar es Salaam 17
Mara       17
Ruvuma     16
Manyara    13
Kigoma     10
Singida    8
Rukwa      2
Name: region, dtype: int64
```

Observations:

- The dataset target variable now has 2 distinct elements.
- Being a vital predictor, age is missing several rows of construction year data and this reduced the size of the dataset to 38,672 rows and 28 columns

NOTE: As evident from the counts, the missing construction years are mainly found in the following areas i.e. Shinyanga(4816), Mbeya(4639), Kagera(3315), Mwanza(2714), Tabora(1959), Dodoma(1840), Iringa(372), Pwani(176) and Lindi(113). with other all other areas missing less than 100 entries each. Analysis carried out with dataset with available age only.

UNIVARIATE ANALYSIS

For this analysis we were interested in the following

- Target Variable Distribution to determine class imbalance
- Age and population distribution of the dataset

```
In [24]: # Get the value counts of the 'status_group' column
value_counts = aged_data['status_group'].value_counts()


# Create an interactive pie chart
fig = px.pie(
    value_counts,
    values=value_counts.values,
    names=value_counts.index,
    # color_discrete_sequence=["cyan", "pink", "yellow"],
    title='Distribution of Status Group',
    hole=0.3
)

display(value_counts)
fig.show()
```

functional	24225
non functional	14447

Name: status_group, dtype: int64

Distribution of Status Group



```

In [25]: # Create subplots with 3 columns
fig, (ax1, ax2, ax3) = plt.subplots(figsize=(16, 6), ncols=3)

# Plot the density of age on the first subplot
sns.kdeplot(aged_data['age'], shade=True, ax=ax1)
ax1.set_title('Distribution of Waterpoint Age')
ax1.set_xlabel('Age')
ax1.set_ylabel('Density')

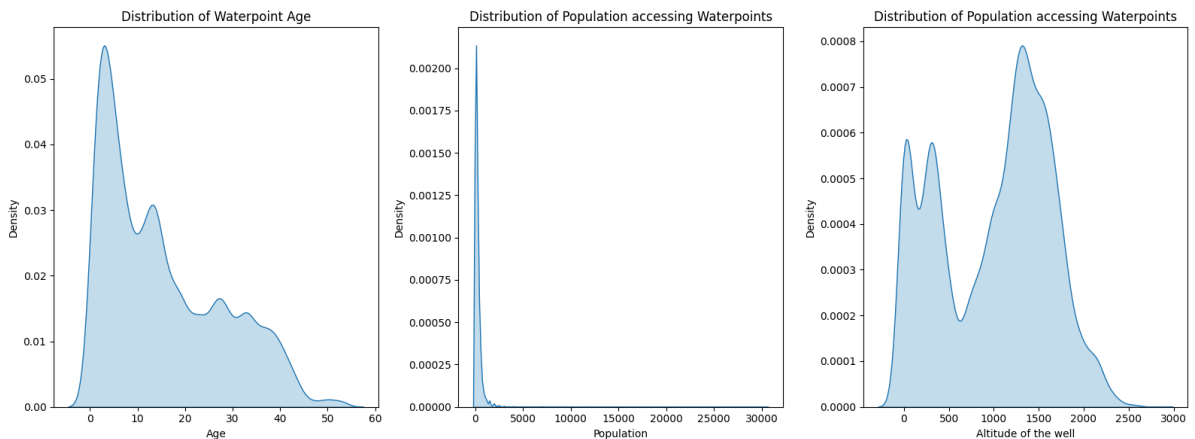
# Plot the density of population on the second subplot
sns.kdeplot(aged_data['population'], shade=True, ax=ax2)
ax2.set_title('Distribution of Population accessing Waterpoints')
ax2.set_xlabel('Population')
ax2.set_ylabel('Density')

# Plot the density of Altitude of the Well on the second subplot
sns.kdeplot(aged_data['gps_height'], shade=True, ax=ax3)
ax3.set_title('Distribution of Population accessing Waterpoints')
ax3.set_xlabel('Altitude of the well')
ax3.set_ylabel('Density')

# Adjust layout to prevent overlap
plt.tight_layout()

# Show the plots
plt.show()

```



```

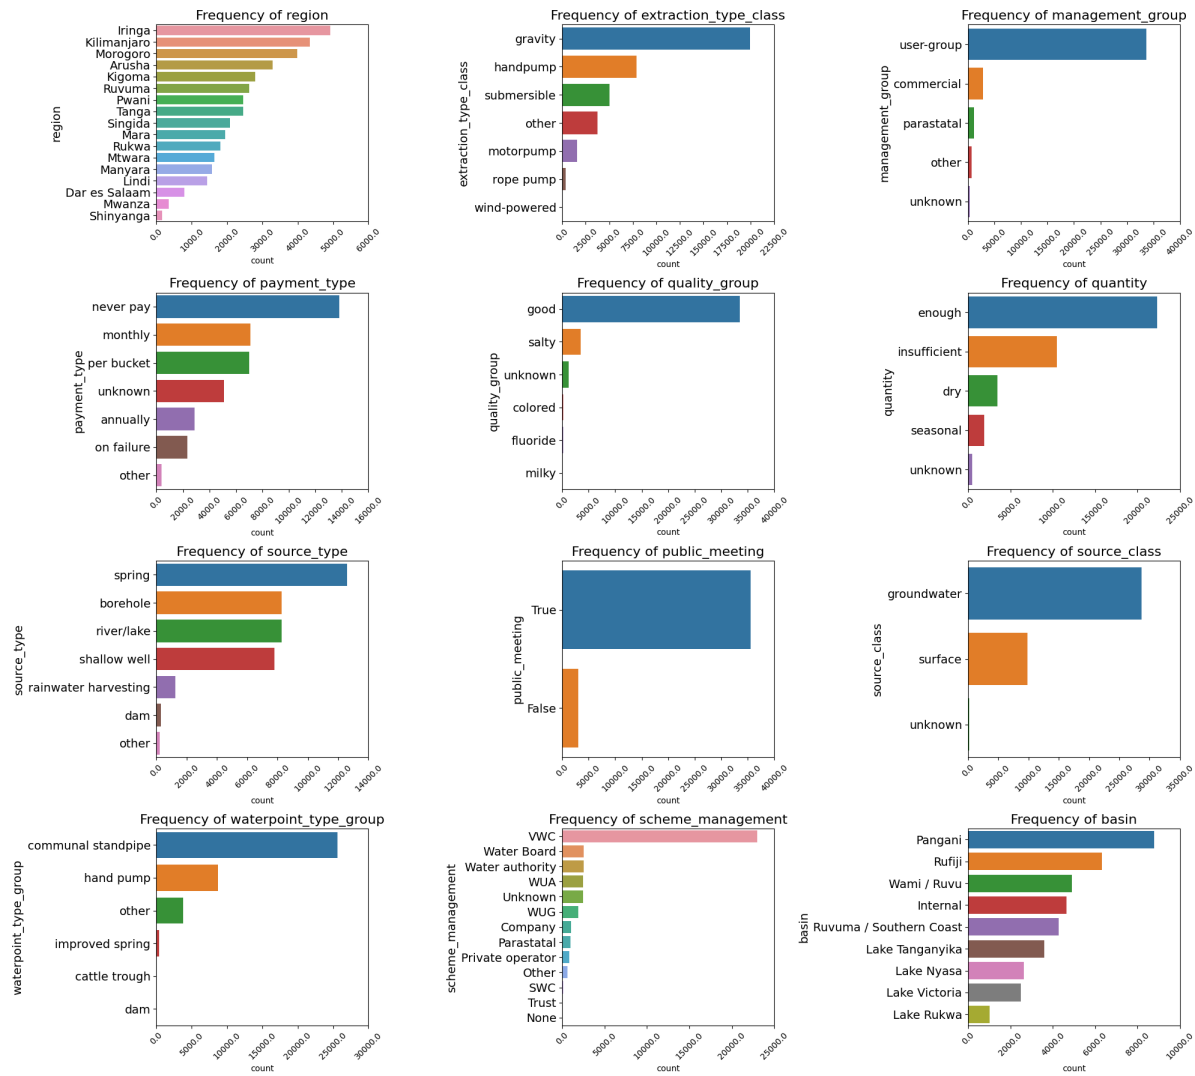
In [26]: def plot_countplots(df, columns):
        """
        Plot count plots for each column in the list
        """
        fig, axes = plt.subplots(4,3, figsize=(20, 18))

        for i in range(len(columns)):
            row = i // 3
            col = i % 3
            ax = axes[row, col]
            sorted_counts = df[columns[i]].value_counts().sort_values(ascending=False)
            sns.countplot(data=df, y=columns[i], order=sorted_counts.index, ax=ax)
            ax.set_title(f'Frequency of {columns[i]}', fontsize=16)
            ax.set_xticks(ax.get_xticks())
            ax.set_xticklabels(ax.get_xticks(), rotation=45, fontsize=10)
            ax.set_yticklabels(ax.get_yticklabels(), fontsize=14)
            ax.set_ylabel(columns[i], fontsize=14)
        plt.tight_layout()
        plt.show()

#List of features to be plotted
columns= ['region', 'extraction_type_class',
          'management_group', 'payment_type',
          'quality_group', 'quantity', 'source_type', 'public_meeting',
          'source_class', 'waterpoint_type_group', 'scheme_management', 'basin'

#plotting countplots
plot_countplots(aged_data, columns, )

```



Observations:

- Our status group show 62.6% functional and 37.4% non functional water points. There is some slight class imbalance

BIVARIATE ANALYSIS

In this section we will consider the relationship between:

- Age and well altitude and the target variable
- Select Features and the target variables
- Deep dive into age, technology, and investment factors relating to water point failure
- Assess socioeconomic and geographic factors relating to water point failure

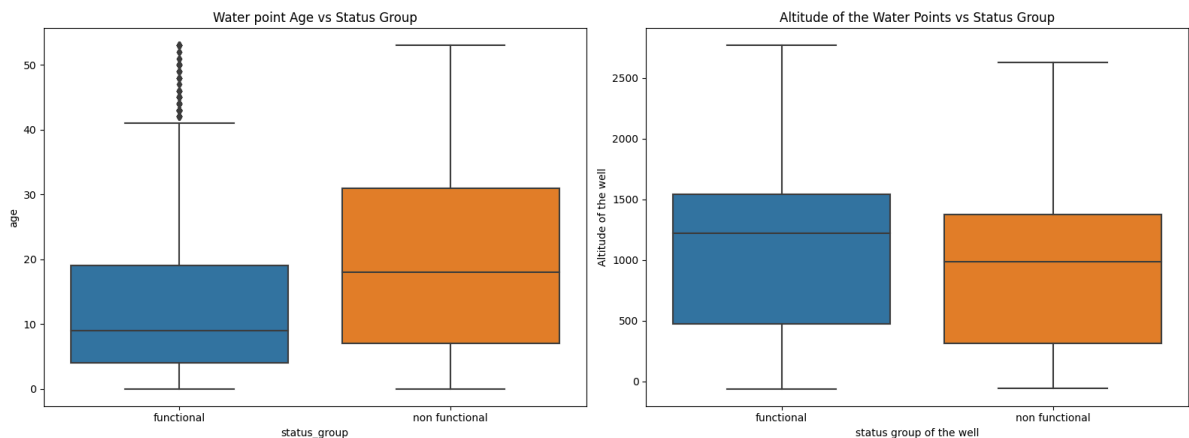
```
In [27]: # Create subplots with 3 columns
fig, (ax1, ax2) = plt.subplots(figsize=(16, 6), ncols=2)

# Plot the density of age
sns.boxplot(x='status_group', y='age', data=aged_data, ax=ax1)
ax1.set_title('Water point Age vs Status Group')
ax1.set_xlabel('status_group')
ax1.set_ylabel('age')

# Plot the density of Altitude of the Well
sns.boxplot(x='status_group', y='gps_height', data=aged_data, ax=ax2)
ax2.set_title('Altitude of the Water Points vs Status Group')
ax2.set_xlabel('status group of the well')
ax2.set_ylabel('Altitude of the well')

# Adjust layout to prevent overlap
plt.tight_layout()

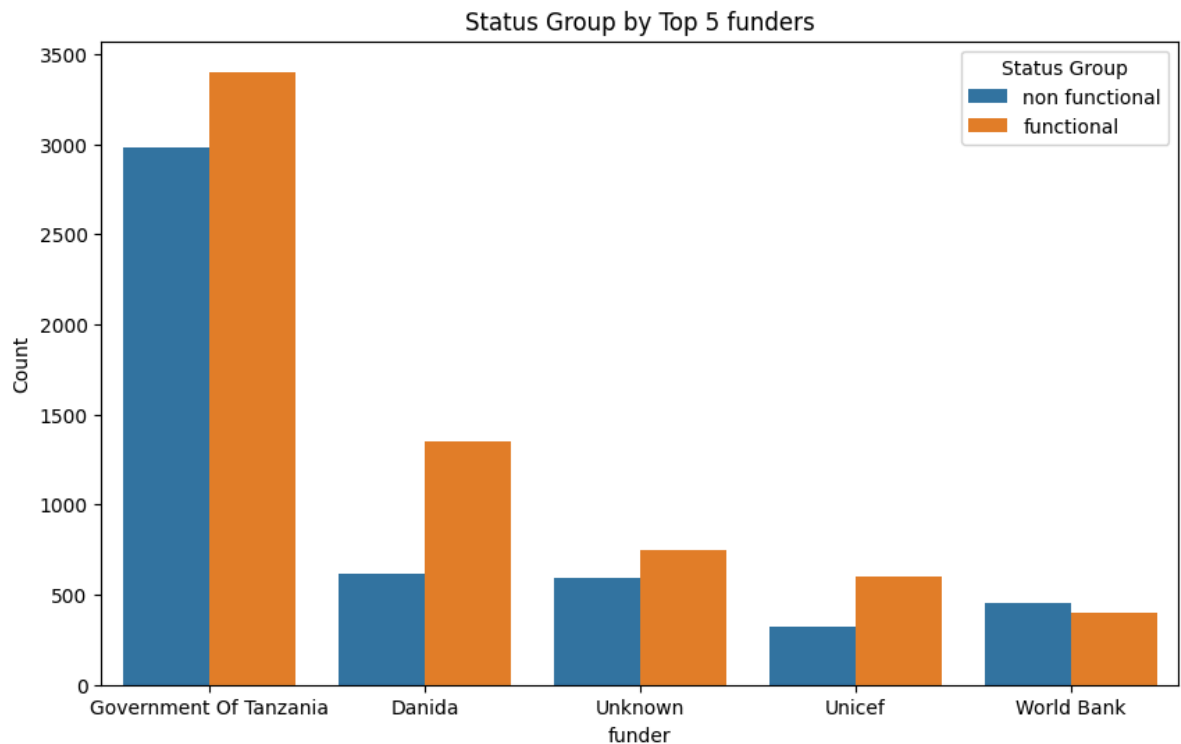
# Show the plots
plt.show()
```

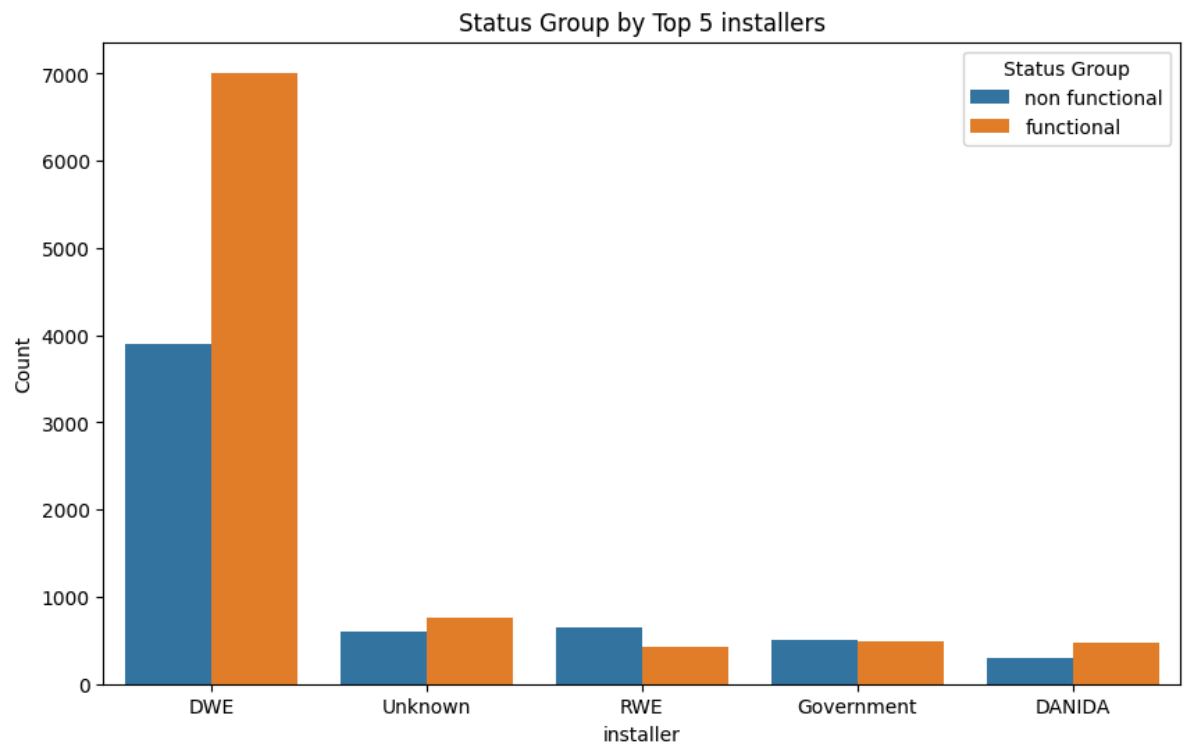



```
In [28]: def plot_count_plots(df, columns):
        """
        Plot count plots for the top five values of each column while keeping hue a
        """
        for column in columns:
            # Get the top n values of the column in descending order
            top_values = df[column].value_counts().head().index

            # Filter the DataFrame to include only the top n values
            df_top_values = df[df[column].isin(top_values)]

            # Create the count plot for the top n values of the column
            plt.figure(figsize=(10, 6))
            sns.countplot(data=df_top_values, x=column, hue="status_group", order=top_values)
            plt.title(f"Status Group by Top 5 {column}s")
            plt.xlabel(column)
            plt.ylabel("Count")
            plt.xticks()
            plt.legend(title="Status Group")
            plt.show()
        plot_count_plots(aged_data, columns=["funder", "installer"])
```





```

In [29]: def plot_count_plots(df, columns, ):
    """
    Plot count plots for each column in the list while keeping hue as "status_g
    """

    num_rows=6
    num_cols=2

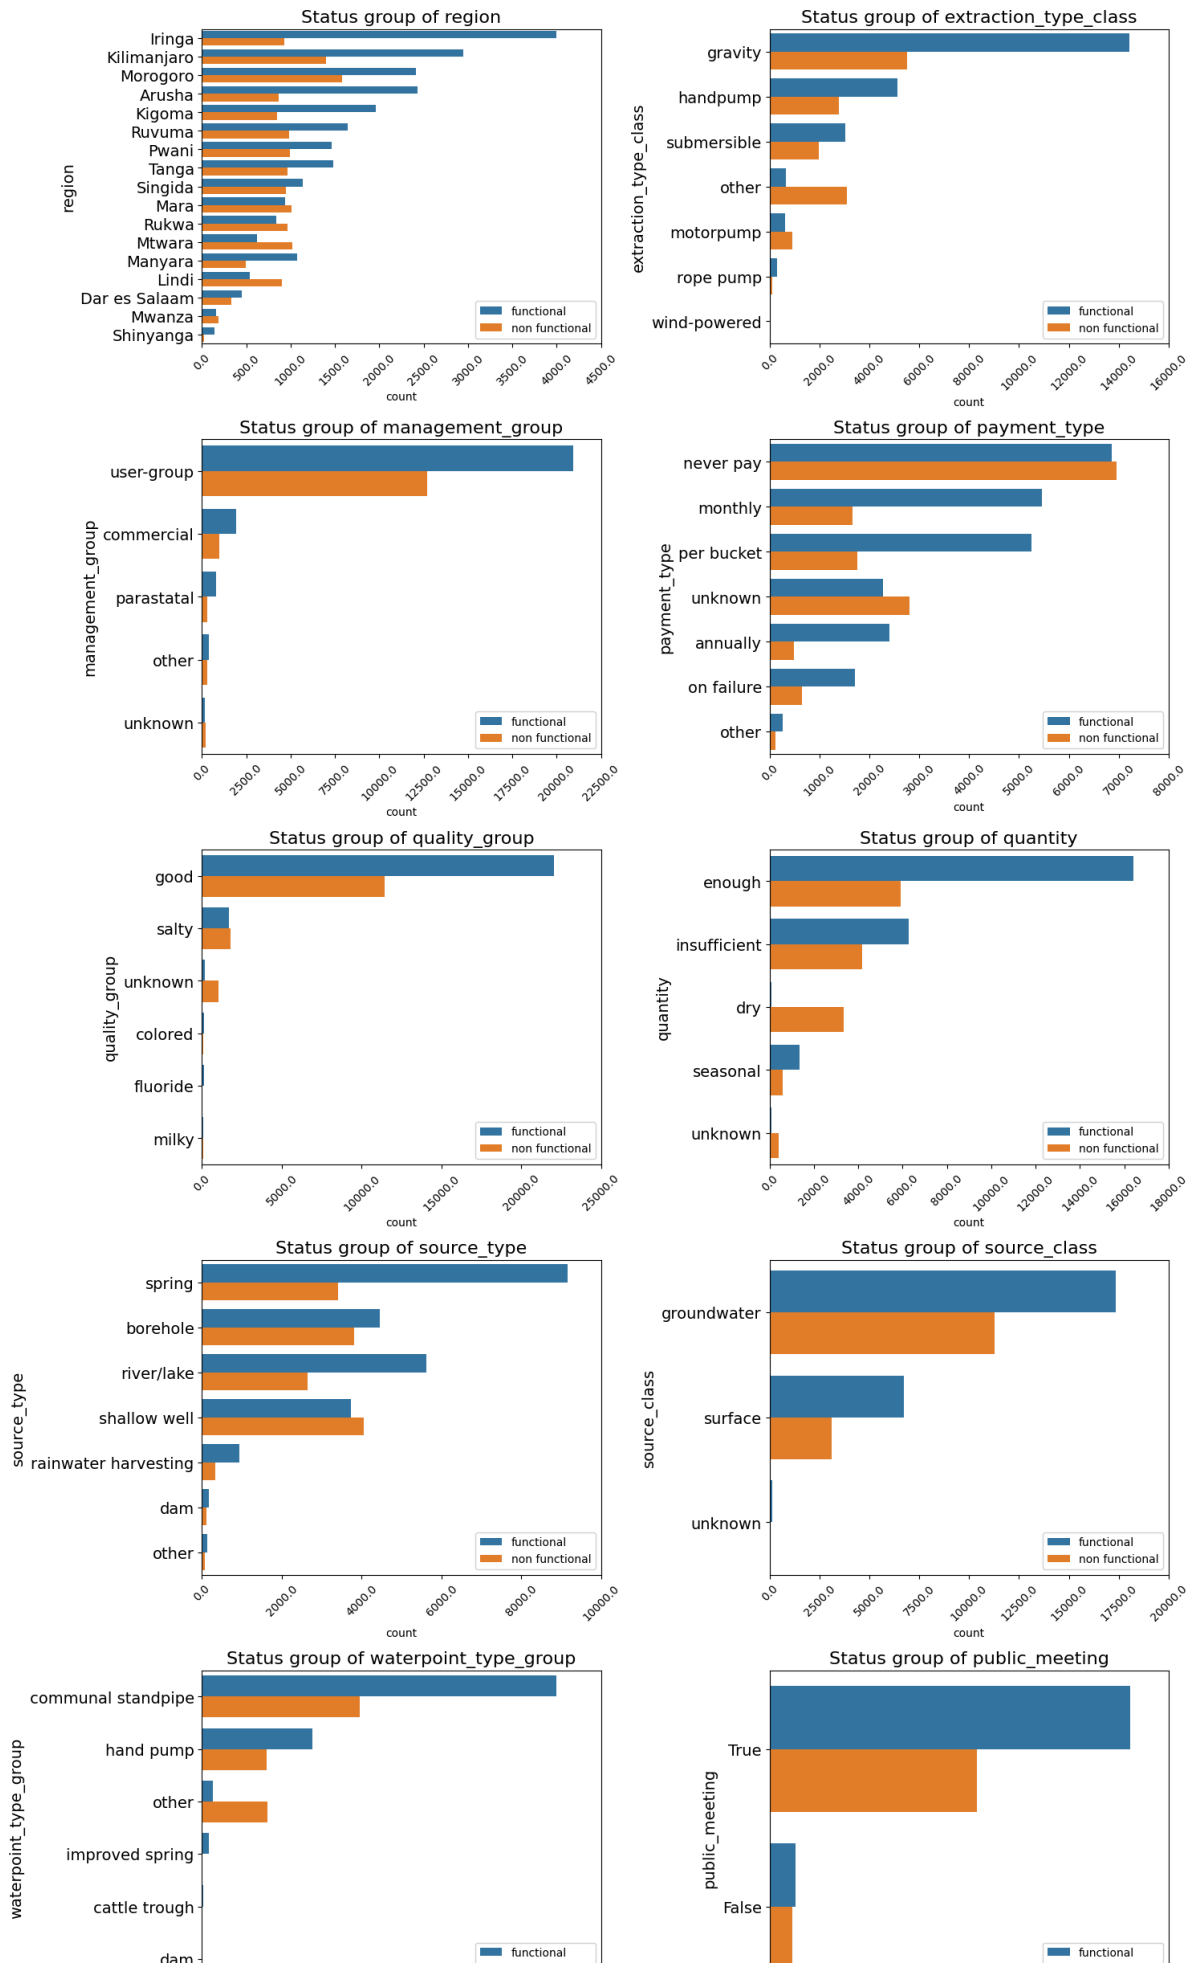
    fig, axes = plt.subplots(num_rows,num_cols, figsize=(15, 30))
    for i in range(len(columns)):
        row = i // 2
        col = i % 2
        ax = axes[row, col]
        sorted_counts = df[columns[i]].value_counts().sort_values(ascending=False)
        sns.countplot(data=df, y=columns[i], order=sorted_counts.index, ax=ax,
        ax.set_title(f'Status group of {columns[i]}', fontsize=16)
        ax.set_xticks(ax.get_xticks())
        ax.set_xticklabels(ax.get_xticks(), rotation=45, fontsize=10)
        ax.set_yticklabels(ax.get_yticklabels(), fontsize=14)
        ax.set_ylabel(columns[i], fontsize=14)
        ax.legend(loc='lower right')


    # removing extra empty axis
    for i in range(len(columns), 6 * 2):
        axes.flatten()[i].remove()

    plt.tight_layout()
    plt.show()

# Columns to be plotted
cols= ['region', 'extraction_type_class',
        'management_group', 'payment_type',
        'quality_group', 'quantity', 'source_type',
        'source_class', 'waterpoint_type_group', "public_meeting"]
plot_count_plots(aged_data, columns=cols )

```



Observations

- We can see that the lifespan of these pumps is roughly 18-20 years with some fairly new pumps failing within the first 10 years
- Waterpoints at a higher altitude generally seem to be fail less as compared to those in lower altitudes
- The local government is the greatest funder of these water points with the distric water engineeer mainly incharge of installing them with most of these waterpoints still in operation
- Key observations from the features are:
 1. Water points where people never pay, have more none functional water points.
 2. Areas with salty water are prone to have more none functional water points
 3. Shallow well source types also seem to have more none functional water points.
 4. Areas with quantity classification as dry have primarily none functional water points.
 5. Motor Pumps equipment are more likely to have none functional water points.
 6. The areas of Mara, Rukwa, Mtwara, Lindi and Mwanza have more none functional waterpoints.

AGE FACTORS

```

In [30]: # Get the value counts of the age below 20 years
value_counts = aged_data[aged_data.age<15]["status_group"].value_counts()

# Create an interactive pie chart
fig = px.pie(
    value_counts,
    values=value_counts.values,
    names=value_counts.index,
    title='Proportion of Water Points Functional In The Last 15 Years',
    hole=0.3
)

display(value_counts)
fig.show()

# Get the value counts of the age ranges
value_counts = aged_data[(aged_data.age>15)&(aged_data.age<25)]["status_group"].value_counts()


# Create an interactive pie chart
fig = px.pie(
    value_counts,
    values=value_counts.values,
    names=value_counts.index,
    title='Proportion of Water Points Functional In Between 15-25 years',
    hole=0.3
)

display(value_counts)
fig.show()

functional      16139
non functional   6050
Name: status_group, dtype: int64


```

Proportion of Water Points Functional In The Last 15 Years



```
functional      3252
non functional  2337
Name: status_group, dtype: int64
```


Proportion of Water Points Functional In Between 15-25 years

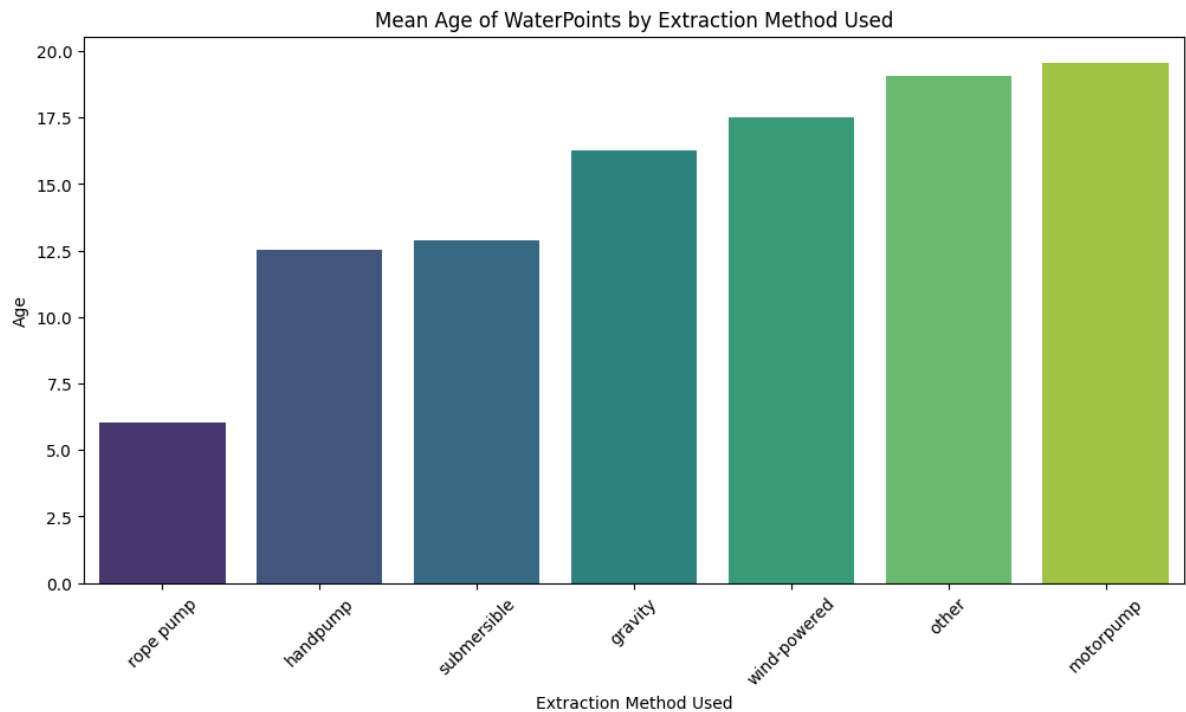


TECHNOLOGICAL FACTORS

```
In [31]: # Group by 'extraction type', calculate the mean 'age', and sort in descending
sorted_data = aged_data.groupby('extraction_type_class')['age'].mean().sort_val

# Convert the Series to a DataFrame
sorted_data_df = sorted_data.reset_index()

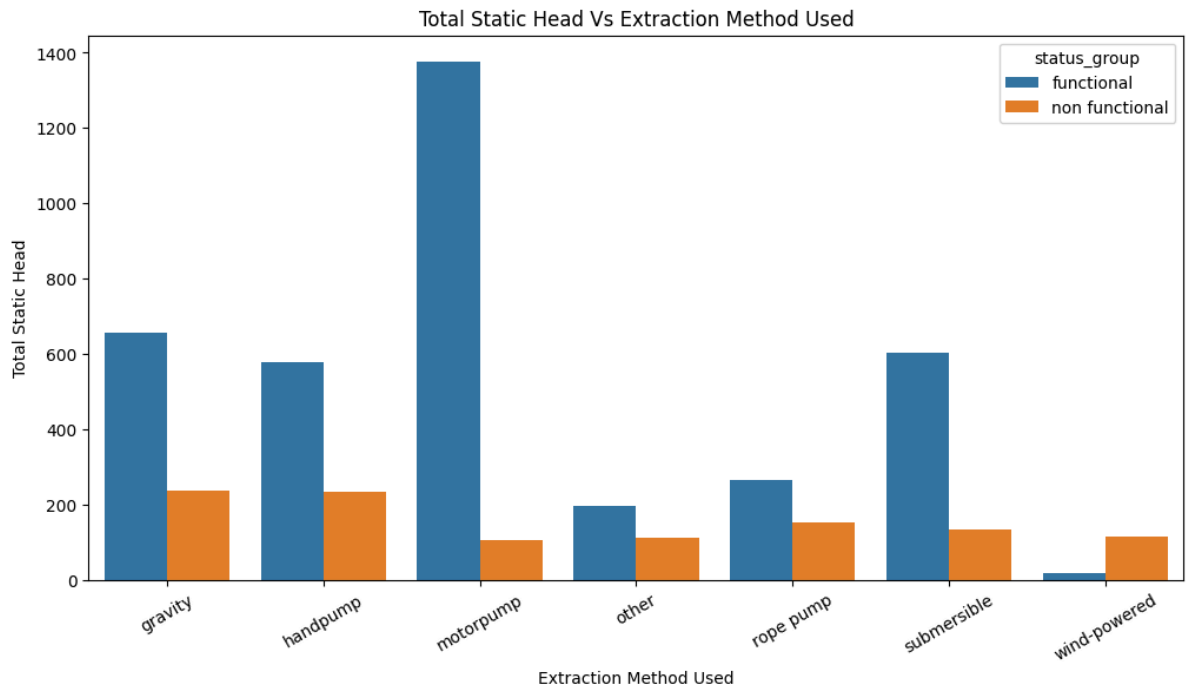
# Plot the data using seaborn
plt.figure(figsize=(12, 6))
sns.barplot(x='extraction_type_class', y='age', data=sorted_data_df, palette='v
plt.title('Mean Age of WaterPoints by Extraction Method Used')
plt.xlabel('Extraction Method Used')
plt.ylabel('Age')
plt.xticks(rotation=45)
plt.show()
```



```
In [32]: # Group by 'extraction type', calculate the mean 'total static head'
sorted_data = aged_data.groupby(['extraction_type_class', 'status_group'])['amount_tsh'].mean()
sorted_data

# Convert the Series to a DataFrame
sorted_data_df = sorted_data.reset_index()

# Plot the data using seaborn
plt.figure(figsize=(12, 6))
sns.barplot(x='extraction_type_class', y='amount_tsh', data=sorted_data_df, hue='status_group')
plt.title('Total Static Head Vs Extraction Method Used')
plt.xlabel('Extraction Method Used')
plt.ylabel('Total Static Head')
plt.xticks(rotation=30)
plt.show()
```

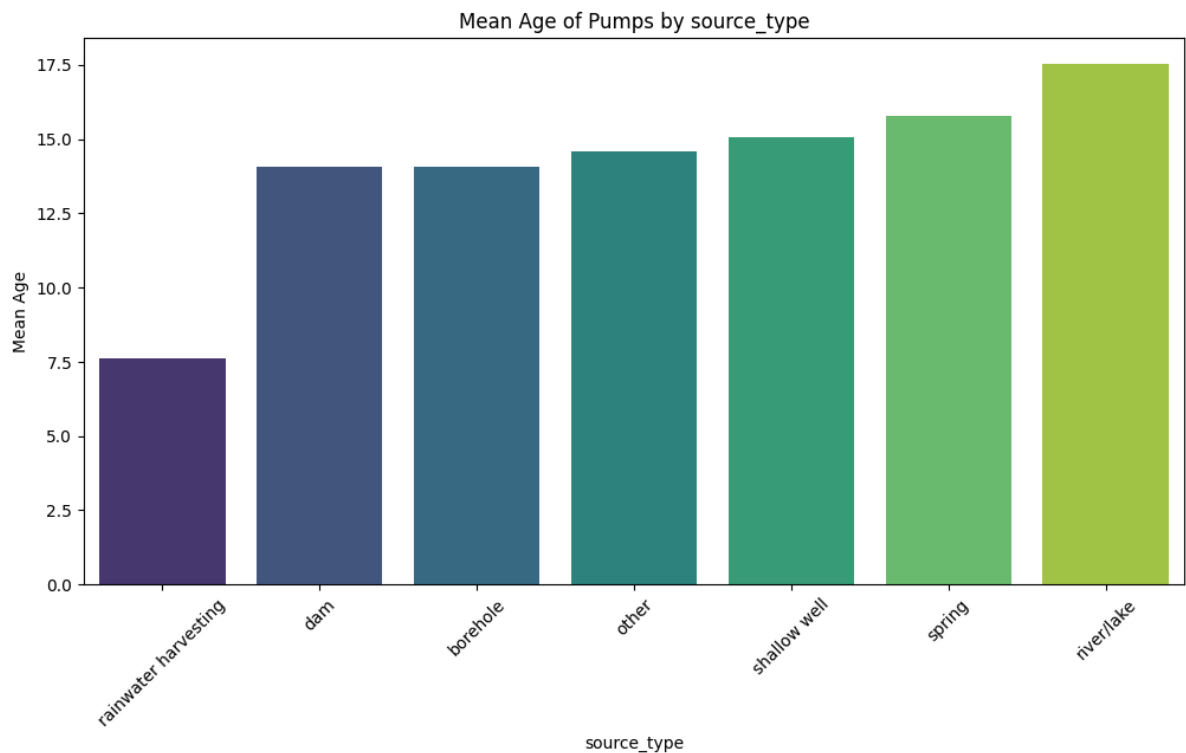


INVESTMENT

```
In [33]: # Group by 'source type', calculate the mean 'age'
sorted_data = aged_data.groupby('source_type')['age'].mean().sort_values(ascending=True)

# Convert the Series to a DataFrame
sorted_data_df = sorted_data.reset_index()

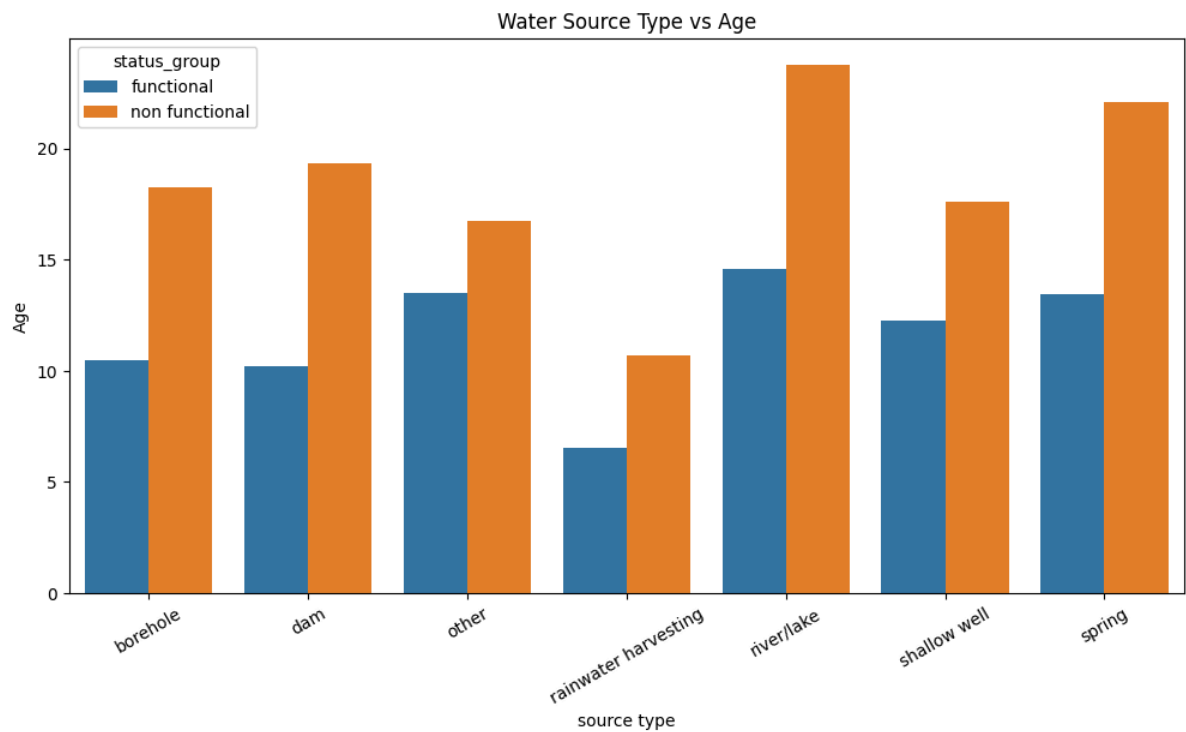
# Plot the data using seaborn
plt.figure(figsize=(12, 6))
sns.barplot(x='source_type', y='age', data=sorted_data_df, palette='viridis')
plt.title('Mean Age of Pumps by source_type')
plt.xlabel('source_type')
plt.ylabel('Mean Age')
plt.xticks(rotation=45)
plt.show()
```



```
In [34]: # Group by 'source type', then by 'status group' calculate the mean 'age'
sorted_data = aged_data.groupby(['source_type', 'status_group'])['age'].mean()
sorted_data

# Convert the Series to a DataFrame
sorted_data_df = sorted_data.reset_index()


# Plot the data using seaborn
plt.figure(figsize=(12, 6))
sns.barplot(x='source_type', y='age', data=sorted_data_df, hue="status_group")
plt.title('Water Source Type vs Age')
plt.xlabel('source type')
plt.ylabel('Age')
plt.xticks(rotation=30)
plt.show()
```



```
In [35]: # Group by 'region', calculate the mean 'age', and sort in descending order
sorted_data = aged_data.groupby('region')['age'].mean().sort_values(ascending=False)

# Convert the Series to a DataFrame
sorted_data_df = sorted_data.reset_index()

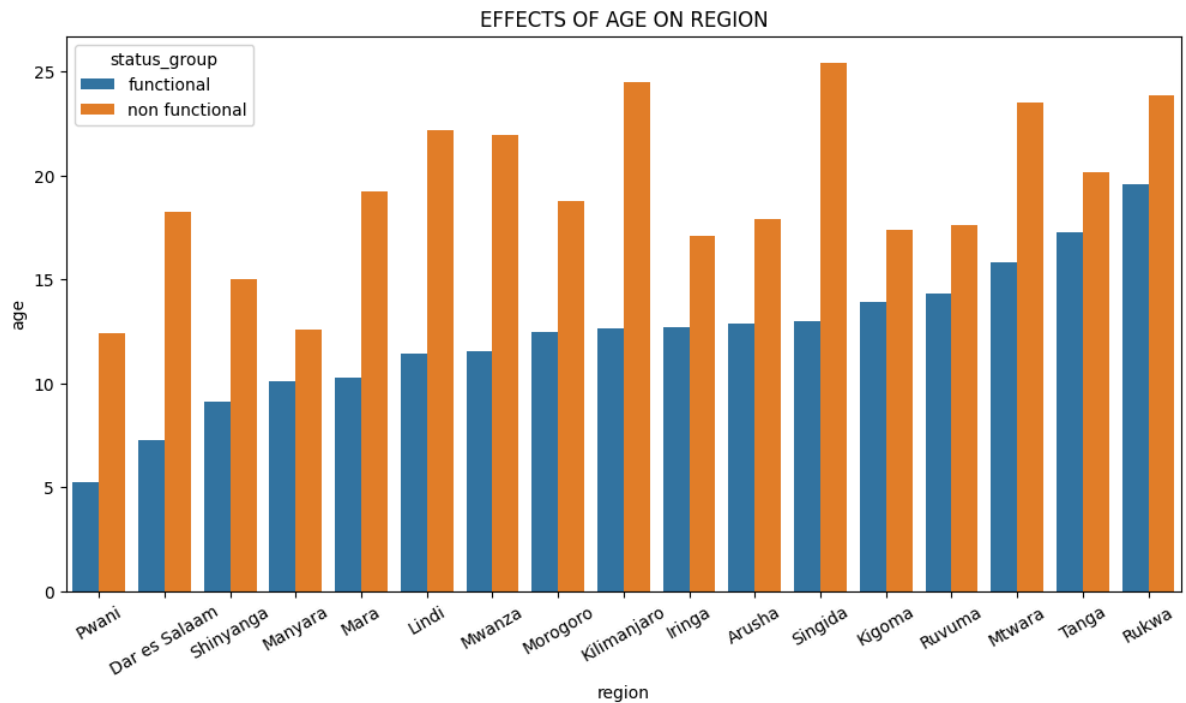
# Plot the data using seaborn
plt.figure(figsize=(12, 6))
sns.barplot(x='region', y='age', data=sorted_data_df, palette='viridis')
plt.title('Mean Age of WaterPoints by Region')
plt.xlabel('extraction_type_class')
plt.ylabel('Age')
plt.xticks(rotation=45)
plt.show()
```



```
In [36]: # Group by 'region', then by 'status group' calculate the mean 'age'
sorted_data = aged_data.groupby(['region', 'status_group'])['age'].mean().sort_
sorted_data

# Convert the Series to a DataFrame
sorted_data_df = sorted_data.reset_index()

# Plot the data using seaborn
plt.figure(figsize=(12, 6))
sns.barplot(x='region', y='age', data=sorted_data_df, hue="status_group")
plt.title('EFFECTS OF AGE ON REGION')
plt.xlabel('region')
plt.ylabel('age')
plt.xticks(rotation=30)
plt.show()
```



FINDINGS AND RESULTS

1. Age Factors

- Majority of the Water Points have been constructed in the last 20 years and are primarily functional i.e. 26,163 water points constructed and 18561 functional.
- On average the life span of these water points is 20 years. After which we see the number of non functional points

2. Technology Factors

- The most common extraction method is gravity type for water access with majority of them functional however we see a majority of motor pumps are non functional.
- Recently there has been a shift to rope pumps and hand pumps as well as submersible pumps. Motor pumps are rarely installed despite motor pumps giving the most amount of water from a waterpoint on average

3. Investment Factors


- Investment has recently been focused on rainwater harvesting, dam and borehole construction in recent years, however it can be seen that despite this investment, majority of these water points are non functional.
- There has been an increased focus around Pwani, Shinyanga, Manyara Regions in construction of water points.

SOCIOECONOMIC FACTORS

```
In [37]: # Group by 'region' then calculate the mean 'population'
sorted_data = aged_data.groupby('region')['population'].mean().sort_values(ascending=True)

# Convert the Series to a DataFrame
sorted_data_df = sorted_data.reset_index()

# Plot the data using seaborn
plt.figure(figsize=(12, 6))
sns.barplot(x='region', y='population', data=sorted_data_df, palette='viridis')
plt.title('Population Access to Water Points By Region')
plt.xlabel('region')
plt.ylabel('population')
plt.xticks(rotation=45)
plt.show()
```



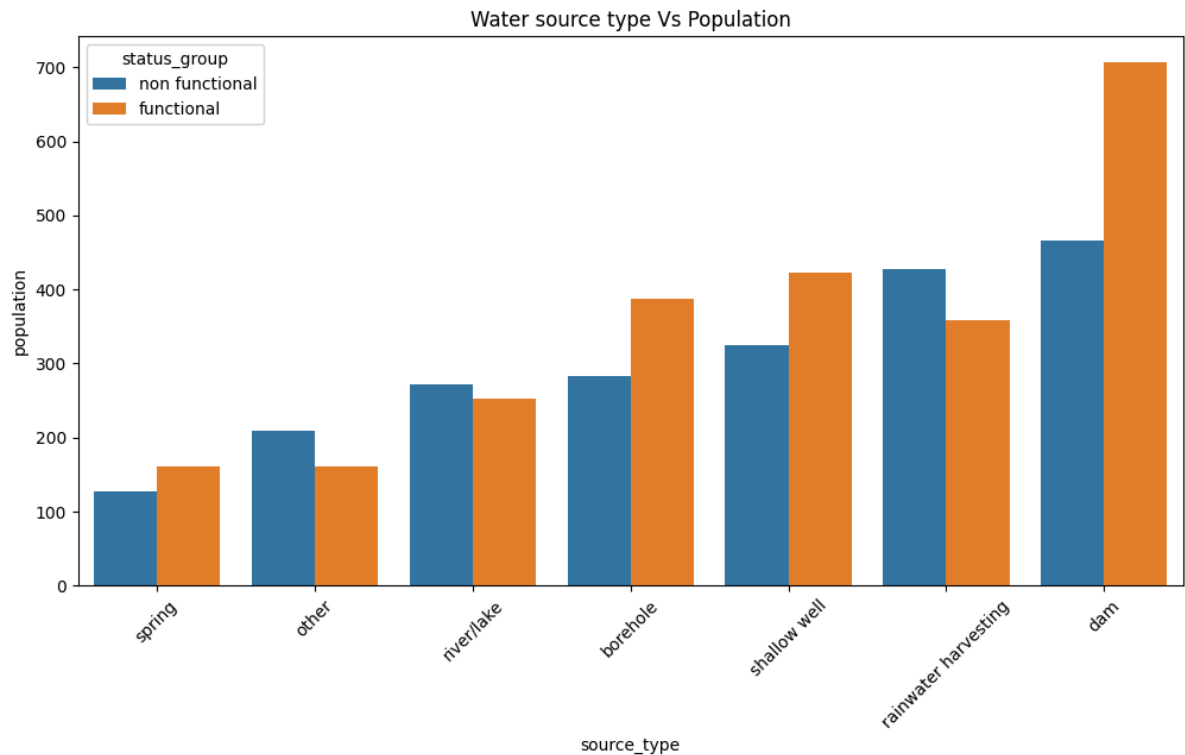
```
In [38]: aged_data.columns
```

```
Out[38]: Index(['amount_tsh', 'date_recorded', 'funder', 'gps_height', 'installer',
               'longitude', 'latitude', 'wpt_name', 'basin', 'subvillage', 'region',
               'lga', 'ward', 'population', 'public_meeting', 'scheme_management',
               'permit', 'construction_year', 'extraction_type_class',
               'management_group', 'payment_type', 'quality_group', 'quantity',
               'source_type', 'source_class', 'waterpoint_type_group', 'status_group',
               'age'],
              dtype='object')
```

```
In [39]: # Group by 'source_type', calculate the mean 'age', and sort in descending order
sorted_data = aged_data.groupby(['source_type', 'status_group'])['population'].sort_values(ascending=False)

# Convert the Series to a DataFrame
sorted_data_df = sorted_data.reset_index()

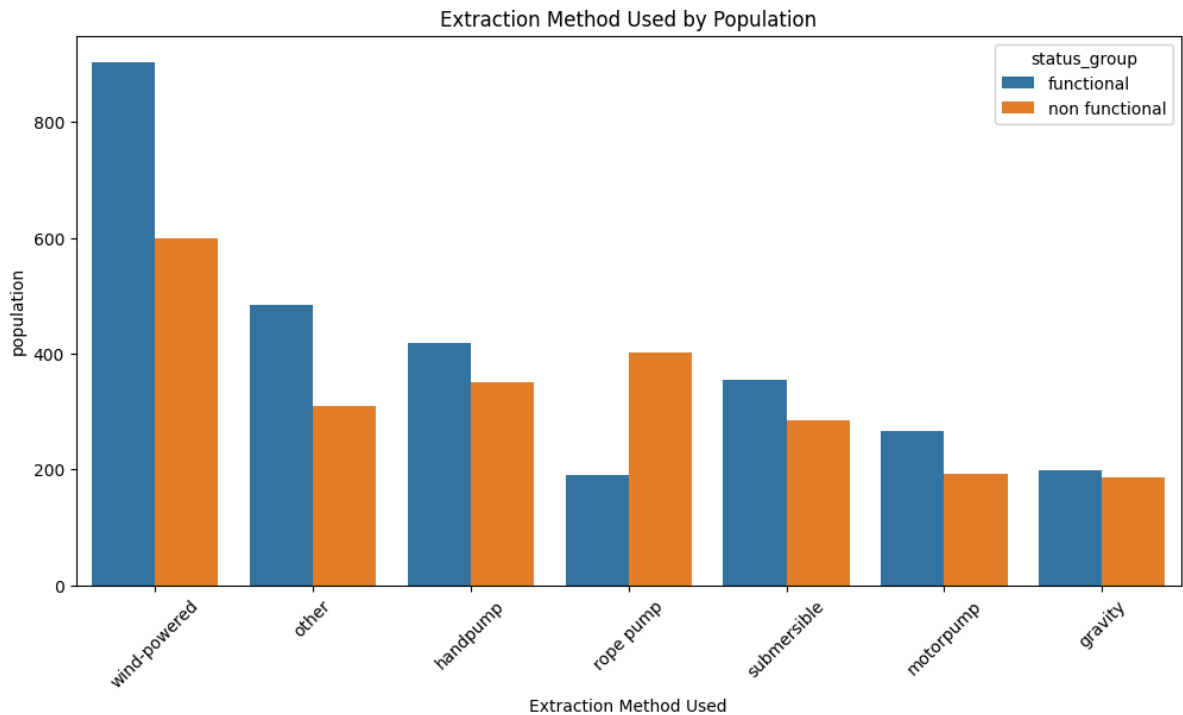
# Plot the data using seaborn
plt.figure(figsize=(12, 6))
sns.barplot(x='source_type', y='population', data=sorted_data_df, hue="status_group")
plt.title('Water source type Vs Population')
plt.xlabel('source_type')
plt.ylabel('population')
plt.xticks(rotation=45)
plt.show()
```



```
In [40]: # Group by 'extraction type', calculate the mean 'age', and sort in descending
sorted_data = aged_data.groupby(['extraction_type_class', 'status_group'])['population'].mean().sort_values(ascending=False)

# Convert the Series to a DataFrame
sorted_data_df = sorted_data.reset_index()

# Plot the data using seaborn
plt.figure(figsize=(12, 6))
sns.barplot(x='extraction_type_class', y='population', data=sorted_data_df, hue='status_group')
plt.title('Extraction Method Used by Population')
plt.xlabel('Extraction Method Used')
plt.ylabel('population')
plt.xticks(rotation=45)
plt.show()
```

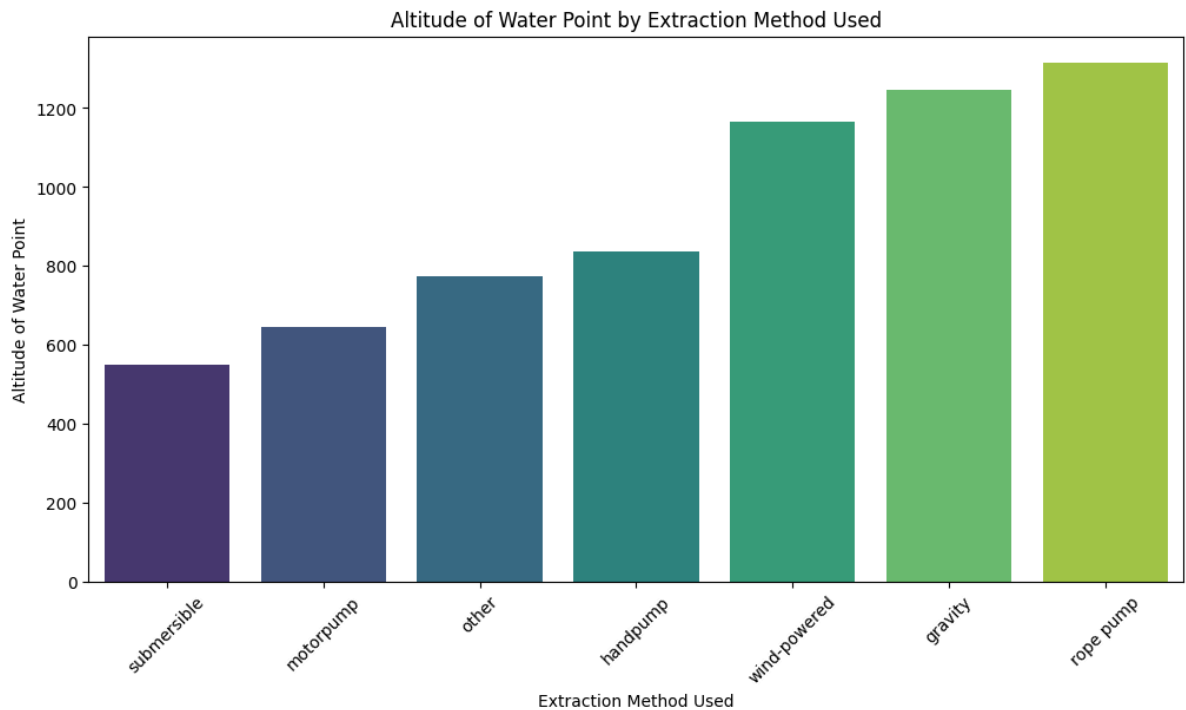


GEOGRAPHICAL FACTORS

```
In [41]: # Group by 'extraction type', calculate the mean 'well altitude', and sort in c
sorted_data = aged_data.groupby('extraction_type_class')['gps_height'].mean().s

# Convert the Series to a DataFrame
sorted_data_df = sorted_data.reset_index()

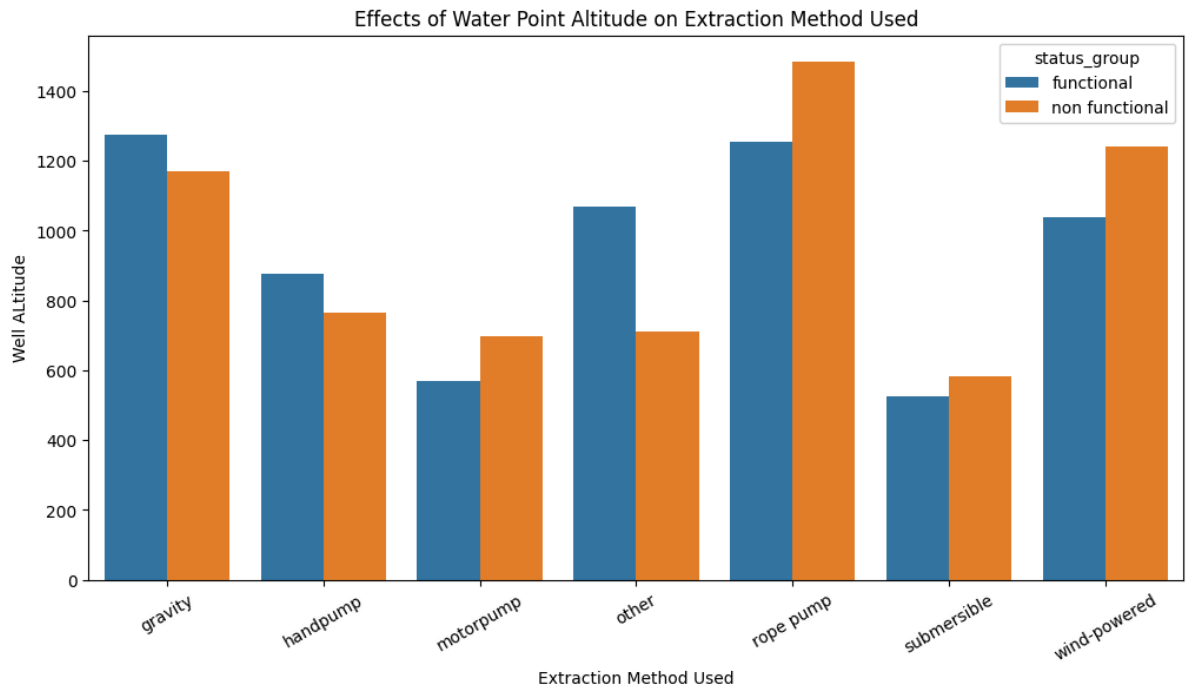
# Plot the data using seaborn
plt.figure(figsize=(12, 6))
sns.barplot(x='extraction_type_class', y='gps_height', data=sorted_data_df, pa
plt.title('Altitude of Water Point by Extraction Method Used')
plt.xlabel('Extraction Method Used')
plt.ylabel('Altitude of Water Point')
plt.xticks(rotation=45)
plt.show()
```



```
In [42]: # Group by 'extraction type', then by 'status group' then calculate the mean 'age'
sorted_data = aged_data.groupby(['extraction_type_class', 'status_group'])['gps_height'].mean()
sorted_data

# Convert the Series to a DataFrame
sorted_data_df = sorted_data.reset_index()

# Plot the data using seaborn
plt.figure(figsize=(12, 6))
sns.barplot(x='extraction_type_class', y='gps_height', data=sorted_data_df, hue='status_group')
plt.title('Effects of Water Point Altitude on Extraction Method Used')
plt.xlabel('Extraction Method Used')
plt.ylabel('Well ALTitude')
plt.xticks(rotation=30)
plt.show()
```



FINDINGS AND RESULTS

1. Social Economic Factors

- Focus needs to be meet in regions such as Lindi, Mwanza, Mara and Rukwa which have high population accessing fewer water points.
- It is also evident that wind powered pumps serve the most population while being less prone to failure

2. Geographic Location Factors

- It is evident that areas with water points in higher altitudes on average have more functional water points.
- The most common technology in higher altitude areas is the rope pump with submersible pumps primarily used in areas of lower altitude.
- We however can see that the rope pumps are prone to failure. Hand Pumps

HANDLING HIGH CARDINALITY COLUMNS

```
In [43]: def high_cardinality_columns(df, threshold=10):
        """
        Display columns with high cardinality based on the threshold.
        """
        high_card_cols = [col for col in df.columns if df[col].nunique() > threshold]

        return print(f"Columns with high cardinality (threshold={threshold}): {high_card_cols}")

high_cardinality_columns(aged_data)
```

Columns with high cardinality (threshold=10): ['amount_tsh', 'funder', 'gps_height', 'installer', 'longitude', 'latitude', 'wpt_name', 'subvillage', 'region', 'lga', 'ward', 'population', 'scheme_management', 'construction_year', 'age']

```
In [44]: # Specific columns chosen to be dropped due to high cardinality and irrelevant
cols=['funder', 'installer', 'longitude', 'latitude',
      'wpt_name', 'subvillage', 'lga', 'ward', 'scheme_management', 'source_class',
      'construction_year', 'date_recorded', 'basin', 'permit', 'public_meeting']
aged_data.drop(columns=cols, inplace=True)
aged_data.shape
```

Out[44]: (38672, 13)

Observations:

- Specific columns with high cardinality were removed as this would cause overfitting in the model
- The dataset has 38,672 rows and 13 rows that shall be used for modelling.

MODELING

The following steps were carried out:

- Data Preprocessing
- Modeling using
 - Logistic Regression Classifier
 - Decision Tree Classifier
 - Random Forest Regression Classifier
 - XGBOOST Classifier
- Hyper Parameter Tuning

DATA PREPROCESSING

A function was created to undertake the following steps for data preprocessing:

- Splitting data
- Onehot encoding the features data
- Label encoding the target data
- Class Imbalance Correction

```

In [45]: def preprocessing(df, test_size=0.25):
        """
        Performs Preprocessing of the data received
        """
        # Splitting the dataset into train and validation sets
        X = df.drop(columns=["status_group"], axis=1)
        y = df["status_group"]

        # Performing train test splits
        X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=test_size)

        # Extract categorical columns
        X_train_cat = X_train.select_dtypes(exclude='number')
        X_val_cat = X_val.select_dtypes(exclude='number')
        X_train_num = X_train.select_dtypes(include='number')
        X_val_num = X_val.select_dtypes(include='number')

        # One-hot encode categorical features
        X_train_enc = pd.get_dummies(X_train_cat)
        X_val_enc = pd.get_dummies(X_val_cat)

        # Instantiate the Label Encoder class
        label_encoder = LabelEncoder()

        # Fit and transform the data
        y_train_lb = label_encoder.fit_transform(y_train)
        y_val_lb = label_encoder.transform(y_val)

        # Return array back into a dataframe
        y_train_enc = pd.Series(y_train_lb, index=y_train.index, name='status_group')
        y_val_enc = pd.Series(y_val_lb, index=y_val.index, name='status_group')

        # Instantiate the SMOTE class
        smote = SMOTE(random_state=42)

        # Fit and resample the data
        X_train_enc, y_train_enc = smote.fit_resample(X_train_enc, y_train_enc)

        return X_train_enc, X_val_enc, y_train_enc, y_val_enc

```

MODELING

A function was created that would undertake classification using different models.


```

In [46]: def modeling(X_train_enc, X_val_enc, y_train_enc, y_val_enc, classifiers):
# Define empty lists to store results in a table
table = []
"""
Inputs different classifiers and return a table with the results for train
"""
# Iterate over classifiers
for clf_name, clf in classifiers.items():
    # Train the model
    clf.fit(X_train_enc, y_train_enc)

    # Make predictions
    y_pred_val = clf.predict(X_val_enc)
    y_pred_train = clf.predict(X_train_enc)

    # Calculate accuracy and F1 score for training set
    train_accuracy = accuracy_score(y_train_enc, y_pred_train)
    train_f1 = f1_score(y_train_enc, y_pred_train)
    train_cross_val = cross_val_score(clf, X_train_enc, y_train_enc, cv=5)

    # Calculate accuracy and F1 score for validation set
    val_accuracy = accuracy_score(y_val_enc, y_pred_val)
    val_f1 = f1_score(y_val_enc, y_pred_val)
    val_cross_val = cross_val_score(clf, X_val_enc, y_val_enc, cv=5).mean()

    # Append results to the list
    table.append({
        'Classifier': clf_name,
        'Data': 'Training',
        'Accuracy': train_accuracy,
        'F1 Score': train_f1,
        'Cross Val Score': train_cross_val
    })
    table.append({
        'Classifier': clf_name,
        'Data': 'Validation',
        'Accuracy': val_accuracy,
        'F1 Score': val_f1,
        'Cross Val Score': val_cross_val
    })

# Create DataFrame from results
model_df = pd.DataFrame(table)

return model_df

```

Prior to modeling, we looked at the best size we should use for the project as below

DETERMINING TEST SIZE

```
In [47]: # List of test sizes to choose from
list_sizes=[0.15,0.2,0.25,0.3]

for size in list_sizes:
    results=[]
    X_train_enc, X_val_enc, y_train_enc,y_val_enc= preprocessing(aged_data, test_data)
    # Define classifiers to build
    classifiers = {
        'Logistic Regression': LogisticRegression(random_state=42),
    }
    # Modeling
    model_results=modeling(X_train_enc, X_val_enc, y_train_enc, y_val_enc, classifiers)
    display(size,model_results)
```

0.15

	Classifier	Data	Accuracy	F1 Score	Cross Val Score
0	Logistic Regression	Training	0.763835	0.741219	0.762455
1	Logistic Regression	Validation	0.777452	0.687636	0.773487

0.2

	Classifier	Data	Accuracy	F1 Score	Cross Val Score
0	Logistic Regression	Training	0.761543	0.738197	0.761002
1	Logistic Regression	Validation	0.775178	0.682026	0.776988

0.25

	Classifier	Data	Accuracy	F1 Score	Cross Val Score
0	Logistic Regression	Training	0.764495	0.741892	0.763371
1	Logistic Regression	Validation	0.775962	0.683888	0.775134

0.3

	Classifier	Data	Accuracy	F1 Score	Cross Val Score
0	Logistic Regression	Training	0.765048	0.741946	0.763787
1	Logistic Regression	Validation	0.774263	0.682507	0.778658

Observation

While using logistic regression with its default hyper parameters we can see that a test size of 0.15 gives us the least variance between the training and test data. Due to the size of the dataset, we do not have enough features to reduce the overfitting issue and will hence use a test size of 0.15 for training to reduce the overfitting

```
In [48]: #Setting the variables with a test size of 0.15
X_train_enc, X_val_enc, y_train_enc, y_val_enc = preprocessing(aged_data, test_size=0.15)

# Define classifiers to build
classifiers = {
    'Logistic Regression': LogisticRegression(random_state=42),
    'KNN': KNeighborsClassifier(),
    'Decision Tree': DecisionTreeClassifier(random_state=42)
}

# Modeling
modeling(X_train_enc, X_val_enc, y_train_enc, y_val_enc, classifiers)
```

Out[48]:

	Classifier	Data	Accuracy	F1 Score	Cross Val Score
0	Logistic Regression	Training	0.763835	0.741219	0.762455
1	Logistic Regression	Validation	0.777452	0.687636	0.773487
2	KNN	Training	0.802287	0.803657	0.787968
3	KNN	Validation	0.779348	0.724257	0.773483
4	Decision Tree	Training	0.834319	0.824532	0.804178
5	Decision Tree	Validation	0.813480	0.743359	0.794348

Observation:

The models were evaluated on their default hyper parameters.

- The best performing model was the decision tree on the accuracy score at 81.34%
- None of the model achieved an f1 score above 75% and so we shall try modeling using ensemble methods

BUILDING MODELS USING ENSEMBLE METHODS

```
In [49]: # Define classifiers to build
classifiers = {
    'Random Forest': RandomForestClassifier(random_state=42),
    'XGBoost': XGBClassifier()
}

# Modeling
modeling(X_train_enc, X_val_enc, y_train_enc, y_val_enc, classifiers)
```

Out[49]:

	Classifier	Data	Accuracy	F1 Score	Cross Val Score
0	Random Forest	Training	0.834270	0.825759	0.808442
1	Random Forest	Validation	0.817618	0.751293	0.803655
2	XGBoost	Training	0.807594	0.793090	0.799308
3	XGBoost	Validation	0.816928	0.745446	0.807102

Observation:

XGBoost though not provide the best scores across the board it has the least variance between train and validation scores.

It was hence choosen for further hyper parameter tuning to try and increase its performance

HYPERPARAMETER TUNING XGBOOST MODEL

To perform some hyper parameter tuning, GridsearchCV was considered but due to the large size of the dataset and the complexity of XGBoost, it was decided to use RandomizedSearchCV for parameter tuning

```
In [50]: # Define the parameter distribution
param_dist = {
    'n_estimators': [100, 200, 500],
    'max_depth': [7, 8, 9],
    'learning_rate': [0.05, 0.1, 0.2],
    'colsample_bytree': [0.3, 0.7, 1.0]
}

# Initialize XGBoost classifier
xgb_clf = XGBClassifier(random_state=42)

# Create the RandomizedSearchCV object
random_search = RandomizedSearchCV(estimator=xgb_clf, param_distributions=param_dist)

# Fit the model
random_search.fit(X_train_enc, y_train_enc)

# Get the best parameters
best_params = random_search.best_params_
print(f"Best parameters found: {best_params}")
```

Fitting 3 folds for each of 81 candidates, totalling 243 fits

Best parameters found: {'n_estimators': 500, 'max_depth': 9, 'learning_rate': 0.2, 'colsample_bytree': 0.3}

```
In [51]: # Assuming best_model is already defined
best_model = random_search.best_estimator_

# Predictions and scores for training data
y_train_pred = best_model.predict(X_train_enc)
f1_train = f1_score(y_train_enc, y_train_pred)
accuracy_train = accuracy_score(y_train_enc, y_train_pred)
cross_val_train= cross_val_score(best_model,y_train_enc, y_train_pred).mean()

# Predictions and scores for validation data
y_val_pred = best_model.predict(X_val_enc)
f1_val = f1_score(y_val_enc, y_val_pred)
accuracy_val = accuracy_score(y_val_enc, y_val_pred)
cross_val_val= cross_val_score(best_model,y_val_enc, y_val_pred).mean()

# Create a DataFrame to store the results
results = pd.DataFrame({
    'Data': ['Training', 'Validation'],
    'F1 Score': [f1_train, f1_val],
    'Accuracy': [accuracy_train, accuracy_val],
    'Cross Validation': [cross_val_train,cross_val_val]
})

# Display the results
results
```

Out[51]:

	Data	F1 Score	Accuracy	Cross Validation
0	Training	0.820346	0.829957	0.829957
1	Validation	0.754307	0.820548	0.820547

```
In [52]: # Assuming rf.feature_importances_ contains the feature importances
feature_importances = best_model.feature_importances_

# Assuming X_train.columns contains the feature labels
feature_labels = X_train_enc.columns

# Create a DataFrame to store feature importances with their corresponding labels
feature_importance_df = pd.DataFrame({'Feature': feature_labels, 'Importance': feature_importances})

# Optionally, you can sort the DataFrame by importance to visualize the most important features
feature_importance_df = feature_importance_df.sort_values(by='Importance', ascending=False)

# Display the DataFrame
feature_importance_df[:5]
```

Out[52]:

	Feature	Importance
42	quantity_dry	0.337651
20	extraction_type_class_other	0.052474
2	region_Iringa	0.050788
58	waterpoint_type_group_improved spring	0.023161
31	payment_type_never pay	0.021514


```
In [53]: # Import confusion_matrix
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

# Print confusion matrix
cnf_matrix = confusion_matrix(y_val_enc, y_val_pred)
print('Confusion Matrix:\n', cnf_matrix)
```

Confusion Matrix:
[[3162 427]
[614 1598]]

```
In [54]: # Visualize your confusion matrix
import matplotlib.pyplot as plt
disp = ConfusionMatrixDisplay(confusion_matrix=cnf_matrix, display_labels=best_
disp.plot(cmap=plt.cm.Blues)
```

```
Out[54]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1d74dd321
60>
```



FINDINGS AND RESULTS

Modeling

- After preprocessing the data and running it through a Logistic Regression and Decision Tree model, the desired accuracy and f1 score was not achieved on either model
- Further modeling out was carried out using a Random Forest Classifier and XGBoost classifier. These results were slightly better but carried some slight overfitting bias to them.
- Using XGBoost classifier the model underwent some parameter tuning and achieved the target accuracy of 80% and target f1 score of 75%

CONCLUSIONS

1. Age of the pumps at the water points are a key indicator for failure with the majority of pumps failing between 15-25 years.
2. Gravity waterpoints give the best performance compared to other forms of technology with motor pumps on the other hand providing the most water per waterpoint i.e. total static head
3. Heavy investment in the past 7-10 years has been on rain water harvesting and has also been focused on the regions of Pwani, Shinyanga and Manyara.
4. Population main source of water are from dams and water harvesting with dams having few failures
5. Areas of high altitude have more functional water points and mainly use rope pumps but lower altitude areas favor submersible pumps

RECOMMENDATIONS

- Water Point Pumps require replacement every 10-15 years to ensure failure doesn't affect the population as well as a premise for predictive maintenance.
- Focus needs to be met in regions such as Lindi, Mwanza, Mara and Rukwa which have high population accessing fewer water points.
- Leverage more reliable extraction type technology such as motor pumps which give more water output per water point. As well as seek to incorporate green technology for sustainability eg solar powered or wind powered technologies
- Using the predictive algorithm, you can predict with upto 80% accuracy to prevent water point downtimes.
- New data required as the dataset is missing significant data points and was recorded over 11 years so feature elements might have changed

In []: